

Operating Systems



4. The Abstraction: The Process

How to provide the illusion of many CPUs?

- ❑ CPU virtualizing

- ◆ The OS can promote the illusion that many virtual CPUs exist.
- ◆ **Time sharing:** Running one process, then stopping it and running another
 - The potential cost is **performance**.

A Process

A process is a **running program**.

- ▣ Comprising of a process:

- ◆ Memory (address space)
 - Instructions
 - Data section
- ◆ Registers
 - Program counter
 - Stack pointer

Process API

- ❑ These APIs are available on any modern OS.
 - ◆ **Create**
 - Create a new process to run a program
 - ◆ **Destroy**
 - Halt a runaway process
 - ◆ **Wait**
 - Wait for a process to stop running
 - ◆ **Miscellaneous Control**
 - Some kind of method to suspend a process and then resume it
 - ◆ **Status**
 - Get some status info about a process


Process Creation

1. Load a program code into memory, into the address space of the process.
 - ◆ Programs initially reside on disk in *executable format*.
 - ◆ OS perform the loading process **lazily**.
 - Loading pieces of code or data only as they are needed during program execution.
2. The program's run-time **stack** is allocated.
 - ◆ Use the stack for *local variables*, *function parameters*, and *return address*.
 - ◆ Initialize the stack with arguments → argc and the argv array of main () function

Process Creation (Cont.)

3. The program's **heap** is created.
 - ◆ Used for explicitly requested dynamically allocated data.
 - ◆ Program request such space by calling `malloc()` and free it by calling `free()`.
4. The OS do some other initialization tasks.
 - ◆ input/output (I/O) setup
 - Each process by default has three open file descriptors.
 - Standard input, output and error
5. **Start the program** running at the entry point, namely `main()`.
 - ◆ The OS *transfers control* of the CPU to the newly-created process.


Loading: From Program To Process



Process States

- A process can be one of three states.
 - ◆ **Running**
 - A process is running on a processor.
 - ◆ **Ready**
 - A process is ready to run but for some reason the OS has chosen not to run it at this given moment.
 - ◆ **Blocked**
 - A process has performed some kind of operation.
 - When a process initiates an I/O request to a disk, it becomes blocked and thus some other process can use the processor.

Process State Transition



Data structures

- ▣ PCB(Process Control Block)
 - ◆ A C-structure that contains information **about each process**.
 - ◆ **Register context:** a set of registers that define the state of a process

- ▣ **Process list**
 - ◆ Ready processes
 - ◆ Blocked processes
 - ◆ Current running processx

Example) thread structure in Pintos

```
struct thread{
    /* Owned by thread.c. */
    tid_t tid;                      /* Thread identifier. */
    enum thread_status status;       /* Thread state. */
    char name[16];                  /* Name (for debugging purposes). */
    uint8_t *stack;                 /* Saved stack pointer. */
    int priority;                   /* Priority. */
    struct list_elem allelem;       /* List element for
all threads list. */

    /* Shared between thread.c and synch.c. */
    struct list_elem elem;          /* List element. */

#ifndef USERPROG
    /* Owned by userprog/process.c. */
    uint32_t *pagedir;              /* Page directory. */
#endif

    /* Owned by thread.c. */
    unsigned magic;                 /* Detects stack overflow. */
};
```

Example) The xv6 kernel Proc Structure (Cont.)

```
// the information xv6 tracks about each process
// including its register context and state
struct proc {
    char *mem;                      // Start of process memory
    uint sz;                         // Size of process memory
    char *kstack;                    // Bottom of kernel stack
                                    // for this process
    enum proc_state state;          // Process state
    int pid;                         // Process ID
    struct proc *parent;            // Parent process
    void *chan;                      // If non-zero, sleeping on chan
    int killed;                      // If non-zero, have been killed
    struct file *ofile[NOFILE];     // Open files
    struct inode *cwd;              // Current directory
    struct context context;          // Switch here to run process
    struct trapframe *tf;           // Trap frame for the
                                    // current interrupt
};
```

Example) Register Context in xv6

```
// the registers xv6 will save and restore
// to stop and subsequently restart a process
struct context {
    int eip;      // Index pointer register
    int esp;      // Stack pointer register
    int ebx;      // Called the base register
    int ecx;      // Called the counter register
    int edx;      // Called the data register
    int esi;      // Source index register
    int edi;      // Destination index register
    int ebp;      // Stack base pointer register
};

// the different states a process can be in
enum proc_state { UNUSED, EMBRYO, SLEEPING,
                  RUNNABLE, RUNNING, ZOMBIE };
```