

Operating Systems



Process API

Overview

- ❑ `fork()`
- ❑ `exec()`
- ❑ `wait()`
- ❑ Separation of `fork()` and `exec()`
 - ◆ IO redirection
 - ◆ pipe

Creating a child process

`fork()`



❑ Create a child process

- ◆ child process is allocated separate memory space from the process. The child process has the same memory contents as the parents.
- ◆ The child process has its own **registers**, and program counter register(**PC**).
- ◆ The newly created process becomes independent after it is created.
- ◆ for parent, `fork()` returns PID of child process; for child process, `fork()` returns 0.


Usage of fork()

p1.c


```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char *argv[]){
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) { // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) { // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else { // parent goes down this path (main)
        printf("hello, I am parent of %d (pid:%d)\n",
               rc, (int) getpid());
    }
    return 0;
}
```

fork



after fork
→



fork(): parent vs. child

Parent vs. Child

```
int pid = fork();  
if (pid > 0) {  
    print ("parent");  
    pid = wait();  
}  
else if (pid == 0) {  
    print ("child");  
    exit();  
}  
else {  
    print ("fork error");  
}
```

parent code

child code

Let's run it.

```
prompt> ./p1
hello world (pid:29146)
hello, I am parent of 29147 (pid:29146)
hello, I am child (pid:29147)
prompt>
```

or

```
prompt> ./p1
hello world (pid:29146)
hello, I am child (pid:29147)
hello, I am parent of 29147 (pid:29146)
prompt>
```

Create the dependency between bewteen the processes

`wait()`

- ❑ When the child process is created, `wait()` in the parent process won't return until the child has run and exited.
- ❑ The parent and the child does not have any dependency.
- ❑ In some cases, the application wants to enforce the order in which they are executed, e.g. the parent exits only after the child finishes.


The usage of wait() System Call

p2.c

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main(int argc, char *argv[]){
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {           // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {    // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else {                // parent goes down this path (main)
        int wc = wait(NULL);
        printf("hello, I am parent of %d (wc:%d) (pid:%d)\n",
               rc, wc, (int) getpid());
    }
    return 0;
}
```

`wait();`



The wait() System Call (Cont.)

Result (Deterministic)

```
prompt> ./p2
hello world (pid:29266)
hello, I am child (pid:29267)
hello, I am parent of 29267 (wc:29267) (pid:29266)
prompt>
```

Running a new program

exec()

- The caller wants to run a program that is different from the caller itself.
 - ◆ Launch an editor
 - ◆ % ls -l
- OS needs to load a new binary image, initialize a new stack, initialize a new heap for the new program.
- two parameters
 - ◆ The name of the binary file
 - ◆ The array of arguments

```
char *argv[3];  
  
argv[0] = "echo";  
argv[1] = "hello";  
argv[2] = 0;  
exec("/bin/echo", argv);  
printf("exec error\n");
```

Usage of exec()


p3.c

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <sys/wait.h>

int main(int argc, char *argv[]){
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) { // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) { // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
        char *myargs[3];
        myargs[0] = strdup("wc"); // program: "wc" (word count)
        myargs[1] = strdup("p3.c"); // argument: file to count
        myargs[2] = NULL; // marks end of array
        execvp(myargs[0], myargs); // runs word count
        printf("this shouldn't print out");
    } else { // parent goes down this path (main)
        int wc = wait(NULL);
        printf("hello, I am parent of %d (wc:%d) (pid:%d)\n",
               rc, wc, (int) getpid());
    }
    return 0;
}
```

When exec() is called,...

- Replace the existing contents of the memory with the new memory contents from the new binary file.
- exec() does not return. It starts to execute the new program.



Usage of exec()

Result

```
prompt> ./p3
hello world (pid:29383)
hello, I am child (pid:29384)
29 107 1030 p3.c
hello, I am parent of 29384 (wc:29384) (pid:29383)
prompt>
```

Why separating `fork()` and `exec()`?

- Why don't we just use something like “`forkandexec("ls", "ls -l")`”?
- Via separating `fork()` and `exec()`, we can manipulate various settings just before executing a new program and **make the IO redirection and pipe possible.**
 - ◆ IO redirection



```
% cat w3.c > newfile.txt
```

- ◆ pipe

```
% echo hello world | wc
```

‘pipe’ is the heart of the shell programming.

IO redirection

```
% wc w3.c > newfile.txt
```

- Save the result of 'wc w3.c' to newfile.txt.
- How?
- Shell is a program that `fork()` and `exec()` the command with argument.
 - ◆ % ls -l → shell calls `fork()` and `exec("ls", "ls-l")` ;
 - ◆ Before calling `exec("wc", "wc w3.c")`, the shell closes STDOUT (`close(1)`) and opens newfile.txt (`open("newfile.txt")`).

Details of IO redirection

p4.c

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <fcntl.h>
#include <sys/wait.h>

int
main(int argc, char *argv[]){
    int rc = fork();
    if (rc < 0) { // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) { // child: redirect standard output to a file
        close(STDOUT_FILENO);
        open("./p4.output", O_CREAT|O_WRONLY|O_TRUNC, S_IRWXU);
// now exec "wc"...
        char *myargs[3];
        myargs[0] = strdup("wc");           // program: "wc" (word count)
        myargs[1] = strdup("p4.c");         // argument: file to count
        myargs[2] = NULL;                  // marks end of array
        execvp(myargs[0], myargs);         // runs word count
    } else {
        int wc = wait(NULL);
    }
    return 0;
}
```

IO redirection


Result

```
prompt> ./p4
prompt> cat p4.output
32 109 846 p4.c
prompt>
```


File descriptor and file descriptor table

File descriptor


- an integer that represents a file, a pipe, a directory and a device
- A process uses a file descriptor to open a file and directory.
- each process has its own file descriptor table.
- File descriptor 0 (Standard Input), 1 (Standard Output), 2 (Standard Error)



File offset



File structure (struct file in Linux)




file descriptor and system calls

- `open()`
 - ◆ Allocate a new file object, allocate new file descriptor and set the newly allocated file descriptor or to point to the new file object.
 - ◆ When allocating the new file descriptor, it uses the smallest ‘free’ file descriptor from the file descriptor table.
- `close()`
 - ◆ deallocate the file descriptor ‘fd’.
 - ◆ Deallocate the file object if there is no file descriptor associated with it.
- `fork()`
 - ◆ copies the file descriptor table from the parent to child process.
- `exec()`
 - ◆ retains the file descriptor table.




fork() and file descriptors

fork and file descriptor




fork()



```
if(fork() == 0) {  
    write(1, "hello ", 6);  
    exit();  
} else {  
    wait();  
    write(1, "world\n", 6);  
}
```

I/O Redirection




file descriptor table



`close(0);
open("input.txt");`

redirection




file descriptor table

`read(0, buf, 1024);`

read 1024 byte from

"input.txt"


I/O Redirection



file descriptor table

`close();
open("output.txt");`

redirection



file descriptor table

`write(1, buf, 1024);`
write the content of
buf to "output.txt".

cat and IO redirection

```
% cat input.txt
```

```
char *argv[2];  
  
argv[0] = "cat";  
argv[1] = 0;  
if(fork() == 0) {  
    close(0);  
    open("input.txt", O_RDONLY);  
    exec("cat", argv);  
}
```


```
char buf[512];  
int n;  
  
for(;;) {  
    n = read(0, buf, sizeof(buf));  
    if(n == 0)  
        break;  
    if(n < 0) {  
        Standard Error  
        fprintf(2, "read error\n");  
        exit();  
    }  
    Standard Output  
    if(write(1, buf, n) != n) {  
        fprintf(2, "write error\n");  
        exit();  
    }  
}
```

pipe: ' | '

```
% echo hello world | wc
```

pipe


- Output to STDOUT of one process is fed to STDIN of another process.
- Implemented with dup() and pipe().
- Key innovation of UNIX shell.



dup (fd)

- duplicate file descriptor: dup() system call

dup (n); // find an empty slot from the beginning
of the file descriptor table.




```
fd = dup(1);
write(1, "hello ", 6);
write(fd, "world\n", 6);
```

pipe()

- special type of file, a kernel buffer that is exposed to a process via a pair of file descriptors: p[0] for read end and p[1] for write end.
- The reader blocks when there is no data to read.

int p[2];

pipe(p);




pipe

```
int p[2];
char *argv[2];

argv[0] = "wc";
argv[1] = 0;

pipe(p);
if(fork() == 0) {
    close(0);
    dup(p[0]);
    close(p[0]);
    close(p[1]);
    exec("/bin/wc", argv);
} else {
    close(p[0]);
    write(p[1], "hello world\n", 12);
    close(p[1]);
}
```




```
% echo hello world | wc
```

pipe

```
int p[2];
char *argv[2];

argv[0] = "wc";
argv[1] = 0;

pipe(p);
if(fork() == 0) {
    close(0);
    dup(p[0]);
    close(p[0]);
    close(p[1]);
    exec("/bin/wc", argv);
} else {
    close(p[0]);
    write(p[1], "hello world\n", 12);
    close(p[1]);
}
```




```
% echo hello world | wc
```

pipe

```
int p[2];
char *argv[2];

argv[0] = "wc";
argv[1] = 0;

pipe(p);
if(fork() == 0) {
    close(0);
    dup(p[0]);
    close(p[0]);
    close(p[1]);
    exec("/bin/wc", argv);
} else {
    close(p[0]);
    write(p[1], "hello world\n", 12);
    close(p[1]);
}
```



```
close(0);
dup(p[0]);
close(p[0]);
```


```
% echo hello world | wc
```

pipe

```
int p[2];
char *argv[2];

argv[0] = "wc";
argv[1] = 0;

pipe(p);
if(fork() == 0) {
    close(0);
    dup(p[0]);
    close(p[0]);
    close(p[1]);
    exec("/bin/wc", argv);
} else {
    close(p[0]);
    write(p[1], "hello world\n", 12);
    close(p[1]);
}
```



pipe vs. IO redirection

- advantages of pipes over using redirection with temporary files

```
echo hello world | wc
```

vs.

```
echo hello world > /tmp/xyz ; wc </tmp/xyz
```

- pipe automatically clean themselves up. When using temporary file, the user has to explicitly delete it.
- pipe can pass arbitrarily long data while file redirection requires sufficient available disk space.
- In pipe, reader and write can proceed in parallel while in redirection, the one has to finish for the others to start.

summary

- ❑ process
- ❑ Process API
 - ◆ fork()
 - ◆ exec()
 - ◆ wait()
- ❑ separation of fork() and exec()
 - ◆ IO redirection
 - ◆ pipe