

# **Operating Systems**

**KAIST**

## **7. Scheduling: Introduction**

---

# Scheduling: Introduction

- ❑ Workload assumptions:

1. Each job runs for the **same amount of time**.
2. All jobs **arrive** at the same time.
3. All jobs only use the **CPU** (i.e., they perform no I/O).
4. The **run-time** of each job is known.

# Scheduling Metrics


- ❑ Performance metric: Turnaround time
  - ◆ The time at which **the job completes** minus the time at which **the job arrived** in the system.

$$T_{turnaround} = T_{completion} - T_{arrival}$$

- ❑ Another metric is fairness.
  - ◆ Performance and fairness are often at odds in scheduling.

# First In, First Out (FIFO)


- ▣ First Come, First Served (FCFS)
  - ◆ Very simple and easy to implement
- ▣ Example:
  - ◆ A arrived just before B which arrived just before C.
  - ◆ Each job runs for 10 seconds.



$$\text{Average turnaround time} = \frac{10 + 20 + 30}{3} = 20 \text{ sec}$$

# Why FIFO is not that great? – Convoy effect


- Let's relax assumption 1: Each job **no longer** runs for the same amount of time.
- Example:
  - A arrived just before B which arrived just before C.
  - A runs for 100 seconds, B and C run for 10 each.



$$\text{Average turnaround time} = \frac{100 + 110 + 120}{3} = 110 \text{ sec}$$

# Shortest Job First (SJF)


- ▣ Run the shortest job first, then the next shortest, and so on
  - ◆ Non-preemptive scheduler
- ▣ Example:
  - ◆ A arrived just before B which arrived just before C.
  - ◆ A runs for 100 seconds, B and C run for 10 each.



$$\text{Average turnaround time} = \frac{10 + 20 + 120}{3} = 50 \text{ sec}$$

# SJF with Late Arrivals from B and C

- Let's relax assumption 2: Jobs can arrive at any time.
- Example:
  - A arrives at t=0 and needs to run for 100 seconds.
  - B and C arrive at t=10 and each need to run for 10 seconds



$$\text{Average turnaround time} = \frac{100 + (110 - 10) + (120 - 10)}{3} = 103.33 \text{ sec}$$


# Shortest Time-to-Completion First (STCF)

- ▣ Add **preemption** to SJF
  - ◆ Also known as Preemptive Shortest Job First (PSJF)
- ▣ A new job enters the system:
  - ◆ Determine of the remaining jobs and new job
  - ◆ Schedule the job which has the least time left

# Shortest Time-to-Completion First (STCF)

## Example:

- ◆ A arrives at t=0 and needs to run for 100 seconds.
- ◆ B and C arrive at t=10 and each need to run for 10 seconds



$$\text{Average turnaround time} = \frac{(120 - 0) + (20 - 10) + (30 - 10)}{3} = 50 \text{ sec}$$

# New scheduling metric: Response time

- The time from **when the job arrives** to the **first time it is scheduled**.

$$T_{\text{response}} = T_{\text{firstrun}} - T_{\text{arrival}}$$

- ◆ STCF and related disciplines are not particularly good for response time.

How can we build a scheduler that is  
sensitive to response time?

# Round Robin (RR) Scheduling


## ▣ Time slicing Scheduling

- ◆ Run a job for a **time slice** and then switch to the next job in the **run queue** until the jobs are finished.
  - Time slice is sometimes called a scheduling quantum.
- ◆ It repeatedly does so until the jobs are finished.
- ◆ The length of a time slice must be *a multiple of* the timer-interrupt period.


RR is fair, but performs poorly on metrics such as turnaround time

# RR Scheduling Example

- A, B and C arrive at the same time.
- They each wish to run for 5 seconds.



$$T_{average\ response} = \frac{0 + 5 + 10}{3} = 5\ sec$$



$$T_{average\ response} = \frac{0 + 1 + 2}{3} = 1\ sec$$

RR with a time-slice of 1sec (Good for Response Time)

# The length of the time slice is critical.


- ▣ The shorter time slice
  - ◆ Better response time
  - ◆ The cost of context switching will dominate overall performance.
  
- ▣ The longer time slice
  - ◆ Amortize the cost of switching
  - ◆ Worse response time

Deciding on the length of the time slice presents  
a **trade-off** to a system designer


# Incorporating I/O

- ▣ Let's relax assumption 3: All programs perform I/O
- ▣ Example:
  - ◆ A and B need 50ms of CPU time each.
  - ◆ A runs for 10ms and then issues an I/O request
    - I/Os each take 10ms
  - ◆ B simply uses the CPU for 50ms and performs no I/O
  - ◆ The scheduler runs A first, then B after

# Incorporating I/O (Cont.)



Poor Use of Resources



Maximize the  
CPU utilization

Overlap Allows Better Use of Resources

# Incorporating I/O (Cont.)

- ▣ When a job initiates an I/O request.
  - ◆ The job is blocked waiting for I/O completion.
  - ◆ The scheduler should schedule another job on the CPU.
- ▣ When the I/O completes
  - ◆ An interrupt is raised.
  - ◆ The OS moves the process from blocked back to the ready state.