

Operating Systems

KAIST

8: Scheduling: The Multi-Level Feedback Queue

Multi-Level Feedback Queue (MLFQ)

- ▣ A Scheduler that learns from the past to predict the future.
- ▣ Objective:
 - ◆ Optimize **turnaround time** → Run shorter jobs first
 - ◆ Minimize **response time** without *a priori knowledge of job length*.

MLFQ: Basic Rules


- ▣ MLFQ has a number of distinct **queues**.
 - ◆ Each queue is assigned a different priority level.
- ▣ A job that is ready to run is on a single queue.
 - ◆ A job **on a higher queue** is chosen to run.
 - ◆ Use round-robin scheduling among jobs in the same queue

Rule 1: If $\text{Priority}(A) > \text{Priority}(B)$, A runs (B doesn't).
Rule 2: If $\text{Priority}(A) = \text{Priority}(B)$, A & B run in RR.

MLFQ: Basic Rules (Cont.)

- ▣ MLFQ varies the priority of a job based on *its observed behavior*.
- ▣ Example:
 - ◆ A job repeatedly relinquishes the CPU while waiting IOs → Keep its priority high
 - ◆ A job uses the CPU intensively for long periods of time → Reduce its priority.

MLFQ Example



Q7


Q6

Q5



Q3

Q2



MLFQ: How to Change Priority


- MLFQ priority adjustment algorithm:

- ◆ **Rule 3:** When a job enters the system, it is placed at the highest priority
- ◆ **Rule 4a:** If a job uses up an entire time slice while running, its priority is reduced (i.e., it moves down on queue).
- ◆ **Rule 4b:** If a job gives up the CPU before the time slice is up, it stays at the same priority level

In this manner, MLFQ approximates SJF

Example 1: A Single Long-Running Job


- A three-queue scheduler with time slice 10ms



Example 2: Along Came a Short Job

- Assumption:


- ◆ **Job A:** A long-running CPU-intensive job
- ◆ **Job B:** A short-running interactive job (20ms runtime)
- ◆ A has been running for some time, and then B arrives at time T=100.



Example 3: What About I/O?

- Assumption:

- Job A: A long-running CPU-intensive job
- Job B: An interactive job that need the CPU only for 1ms before performing an I/O





The MLFQ approach keeps an interactive job at the highest priority

Problems with the Basic MLFQ

- ▣ Starvation
 - ◆ If there are "too many" interactive jobs in the system.
 - ◆ Long-running jobs will never receive any CPU time.
- ▣ Game the scheduler
 - ◆ After running 99% of a time slice, issue an I/O operation.
 - ◆ The job gain a higher percentage of CPU time.
- ▣ A program may change its behavior over time.
 - ◆ CPU bound process → I/O bound process

The Priority Boost

- **Rule 5:** After some time period S , move all the jobs in the system to the topmost queue.
 - ◆ Example:
 - A long-running job(A) with two short-running interactive job(B, C)





Without(Left) and With(Right) Priority Boost

A: B: C:

Better Accounting

- How to prevent gaming of our scheduler?
- Solution:
 - ◆ **Rule 4** (Rewrite Rules 4a and 4b): Once a job **uses up its time allotment** at a given level (regardless of how many times it has given up the CPU), **its priority is reduced**(i.e., it moves down on queue).




Without(Left) and With(Right) Gaming Tolerance

Tuning MLFQ And Other Issues

Lower Priority, Longer Quanta

- ◆ The high-priority queues → Short time slices
 - E.g., 10 or fewer milliseconds
- ◆ The Low-priority queue → Longer time slices
 - E.g., 100 milliseconds



Example) 10ms for the highest queue, 20ms for the middle,
40ms for the lowest

The Solaris MLFQ implementation

- ▣ For the Time-Sharing scheduling class (TS)
 - ◆ 60 Queues
 - ◆ Slowly increasing time-slice length
 - The highest priority: 20msec
 - The lowest priority: A few hundred milliseconds
 - ◆ Priorities boosted around every 1 second or so.

FreeBSD Scheduler(4.3)

- ▣ MLFQ without queue.
- ▣ Instead, use formula.
- ▣ Compute the priority of a process based upon
 - ◆ How much CPU a process has used.
 - ◆ Boost priority by decay.
 - ◆ Take the advice from the user (nice).
- ▣ For efficiency, use queue.

MLFQ: Summary

- ▣ The refined set of MLFQ rules:
 - ◆ **Rule 1:** If Priority(A) > Priority(B), A runs (B doesn't).
 - ◆ **Rule 2:** If Priority(A) = Priority(B), A & B run in RR.
 - ◆ **Rule 3:** When a job enters the system, it is placed at the highest priority.
 - ◆ **Rule 4:** Once a job uses up its time allotment at a given level (regardless of how many times it has given up the CPU), its priority is reduced(i.e., it moves down on queue).
 - ◆ **Rule 5:** After some time period S, move all the jobs in the system to the topmost queue.
- ▣ Beauty of MLFQ
 - ◆ It does not require prior knowledge on the CPU usage of a process.