

Operating Systems

KAIST

9: Scheduling: Proportional Share

Proportional Share Scheduler

- ▣ Fair-share scheduler

- ◆ Guarantee that each job obtain *a certain percentage* of CPU time.
- ◆ Not optimized for turnaround or response time

Basic Concept

▣ Tickets

- ◆ Represent the share of a resource that a process should receive
- ◆ The percent of tickets represents its share of the system resource in question.

▣ Example

- ◆ There are two processes, A and B.
 - Process A has 75 tickets → receive 75% of the CPU
 - Process B has 25 tickets → receive 25% of the CPU

Lottery scheduling

- ▣ The scheduler picks a winning ticket.
 - ◆ Load the state of that *winning process* and runs it.
- ▣ Example
 - ◆ There are 100 tickets
 - Process A has 75 tickets: 0 ~ 74
 - Process B has 25 tickets: 75 ~ 99

Scheduler's winning tickets: 63 85 70 39 76 17 29 41 36 39 10 99 68 83 63

Resulting scheduler: A B A A B A A A A A A B A B A

The longer these two jobs compete,
The more likely they are to achieve the desired percentages.

Ticket Mechanisms

❑ Ticket currency

- ◆ A user allocates tickets among their own jobs in whatever currency they would like.
- ◆ The system converts the currency into the correct global value.
- ◆ Example
 - There are 200 tickets (Global currency)
 - Process A has 100 tickets
 - Process B has 100 tickets

User A → 500 (A's currency) to A1 → 50 (global currency)
→ 500 (A's currency) to A2 → 50 (global currency)

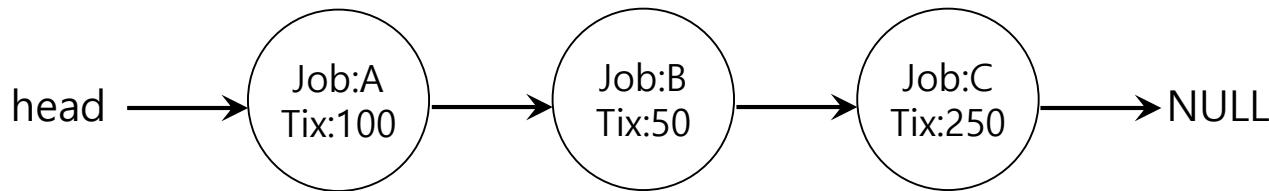
User B → 10 (B's currency) to B1 → 100 (global currency)

Ticket Mechanisms (Cont.)

- ▣ Ticket transfer
 - ◆ A process can temporarily hand off its tickets to another process.
- ▣ Ticket inflation
 - ◆ A process can temporarily raise or lower the number of tickets it owns.
 - ◆ If any one process needs *more CPU time*, it can boost its tickets.

Implementation

- Example: There are three processes, A, B, and C.
 - Keep the processes in a list sorted with the ticket size: highest ticket first



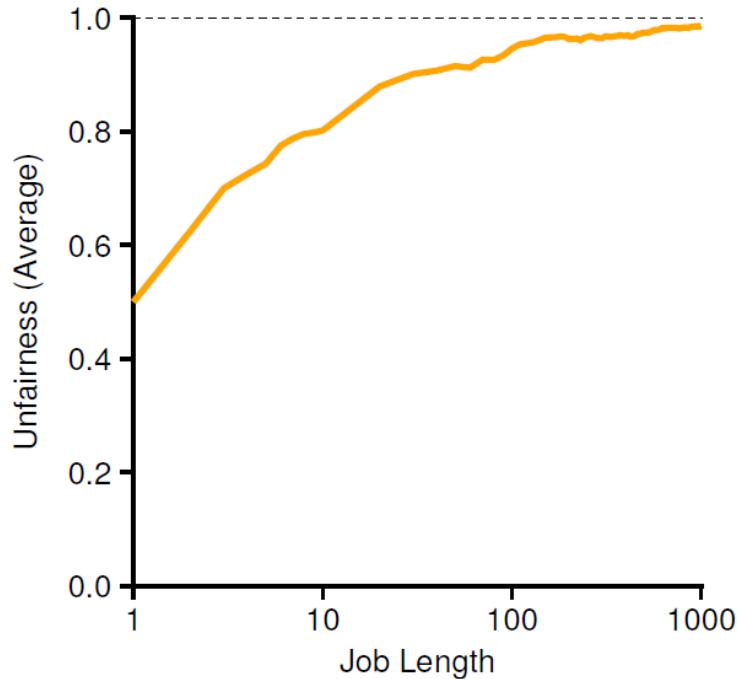
```
1 // counter: used to track if we've found the winner yet
2 int counter = 0;
3
4 // winner: use some call to a random number generator to
5 // get a value, between 0 and the total # of tickets
6 int winner = getrandom(0, totaltickets);
7
8 // current: use this to walk through the list of jobs
9 node_t *current = head;
10
11 // loop until the sum of ticket values is > the winner
12 while (current) {
13     counter = counter + current->tickets;
14     if (counter > winner)
15         break; // found the winner
16     current = current->next;
17 }
18 // 'current' is the winner: schedule it...
```

Implementation (Cont.)

- ▣ U: unfairness metric
 - ◆ The time the first job completes divided by the time that the second job completes.
- ▣ Example:
 - ◆ There are two jobs, each job has runtime 10.
 - First job finishes at time 10
 - Second job finishes at time 20
 - ◆ $U = \frac{10}{20} = 0.5$
 - ◆ U will be close to 1 when both jobs finish at nearly the same time.

Lottery Fairness Study

- There are two jobs.
 - Each job has the same number of tickets (100).



When the job length is not very long,
average unfairness can be quite severe.

Deterministic Approach: Stride Scheduling

- ▣ **Stride** of each process
 - ◆ (A large number) / (the number of tickets of the process)
 - ◆ Example: A large number = 10,000
 - Process A has 100 tickets → stride of A is 100
 - Process B has 50 tickets → stride of B is 200
- ▣ A process runs, increment a counter(=pass value) for it by its stride.
 - ◆ Pick the process to run that has **the lowest pass value**

```
current = remove_min(queue);           // pick client with minimum pass
schedule(current);                   // use resource for quantum
current->pass += current->stride;    // compute next pass using stride
insert(queue, current);              // put back into the queue
```

A pseudo code implementation

Stride Scheduling Example

Pass(A) (stride=100)	Pass(B) (stride=200)	Pass(C) (stride=40)	Who Runs?
0	0	0	A
100	0	0	B
100	200	0	C
100	200	40	C
100	200	80	C
100	200	120	A
200	200	120	C
200	200	160	C
200	200	200	...

Stride scheduling needs to maintain the per process pass value.
If new job enters with pass value 0 it will monopolize the CPU!

Advantage of Lottery scheduling: no per-process state

The Linux Completely Fair Scheduling (CFS)

❑ Completely Fair Scheduling (CFS)

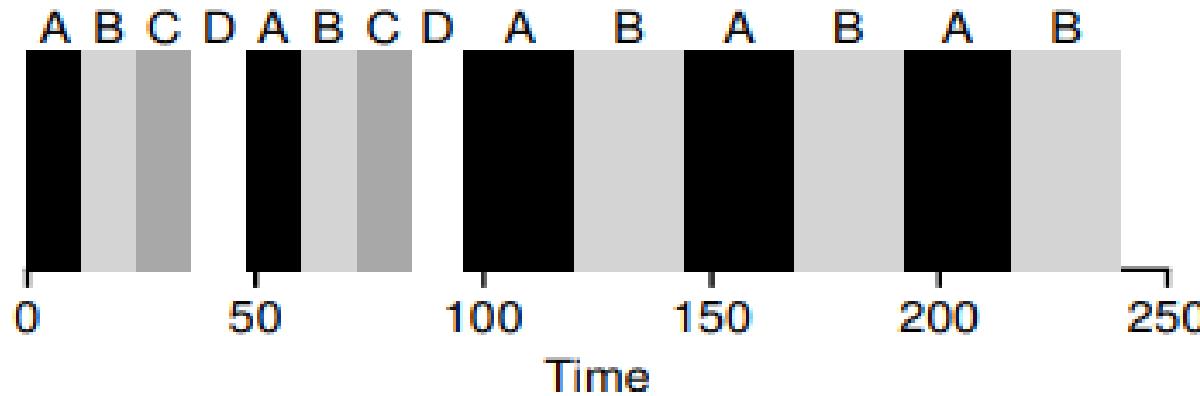
- ◆ The current CPU scheduler in Linux
- ◆ Non-fixed timeslice.
 - CFS assigns process's timeslice a proportion of the processor.
- ◆ Priority
 - Enables control over priority by using nice value.
- ◆ Efficient data structure.
 - Use red-black tree for efficient search, insertion and deletion of a process.

Basic

- ▣ Virtual runtime (vruntime)
 - ◆ Denote how long the process has been executing.
 - ◆ Per-process variable
 - ◆ Increase in **proportion with physical (real) time** when it runs.
 - ◆ CFS will pick the process with the **lowest vruntime** to run next.
- ▣ sched_latency
 - ◆ A typical value is 48 (milliseconds)
 - ◆ process's timeslice = sched_latency / (the number of process)

Example

- ◆ Simple Example
 - 4 processes (A,B,C,D) and then 2 processes(C,D) complete.



- ◆ min_granularity
 - The minimum timeslice (6ms)
 - Ensure that not too much time is spent in scheduling overhead, When there are too many processes running.

Weight

- ◆ Nice value
 - CFS enables controls over process priority.
 - Nice parameter is integer value and can be set from -20 to +19.
 - The nice value is mapped to a weight (value is not important)

```
static const int prio_to_weight[40] = {
    /* -20 */     88761,      71755,      56483,      46273,      36291,
    /* -15 */     29154,      23254,      18705,      14949,      11916,
    /* -10 */     9548,       7620,       6100,       4904,       3906,
    /* -5 */      3121,       2501,       1991,       1586,       1277,
    /*  0 */      1024,        820,        655,        526,        423,
    /*  5 */      335,         272,        215,        172,        137,
    /* 10 */      110,          87,         70,         56,         45,
    /* 15 */      36,           29,         23,         18,         15,
};
```

Weighting (Niceness)

- ◆ New timeslice formula

$$time_slice_k = \frac{weight_k}{\sum_{n=0}^{n-1} weight_i} \cdot sched_latency$$

- ◆ Simple Example
 - Assign Process `A` a nice value of -5 and process `B` a nice value of 0.

Process	nice value	weight	Time slice
A	-5	3121	36 ms
B	0	1024	12 ms

vruntime with weight

- Weighting (Niceness)

- vruntime formula

- Calculate the actual run time. Scales it inversely by the weight of process.

$$vruntime_i = vruntime_i + \frac{weight_0}{weight_i} \cdot runtime_i$$

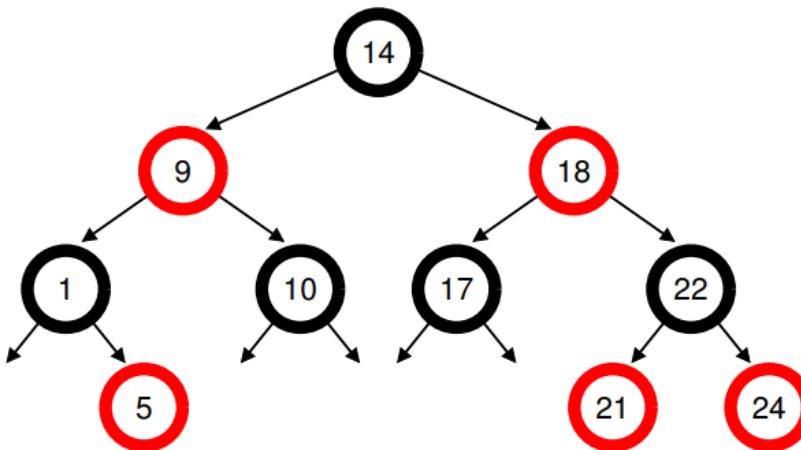
- Simple Example

Process	nice value	weight	Accumulated value
A	-5	3121	1 * runtime
B	0	1024	3 * runtime

Structure of ready queue

▫ Red-Black Tree

- ◆ Balanced binary tree (can address worst-case insertion)
- ◆ Ordering of Red-Black Tree : $O(\log n)$
- ◆ Efficiently find the process with minimum virtual runtime.
- ◆ Only running (or runnable) processes are kept therein.



I/O and sleeping process

- ▣ Dealing with I/O and Sleeping processes
 - ◆ Avoid the situation where some process monopolizes the CPU, if process have significantly small vruntime after sleeping.
 - ◆ Set the vruntime of process to the minimum value found in tree when it wakes up.
 - ◆ Process that sleep for short periods of time frequently do not ever get their fair share of the CPU.