

Operating Systems

KAIST

18. Paging: Introduction

Concept of Paging

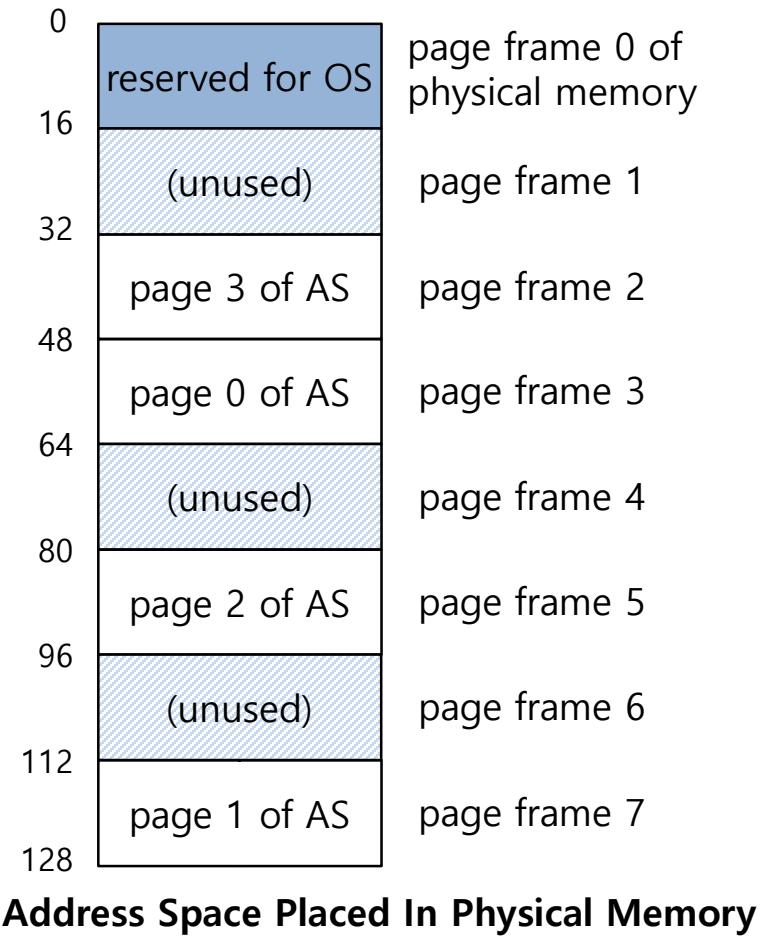
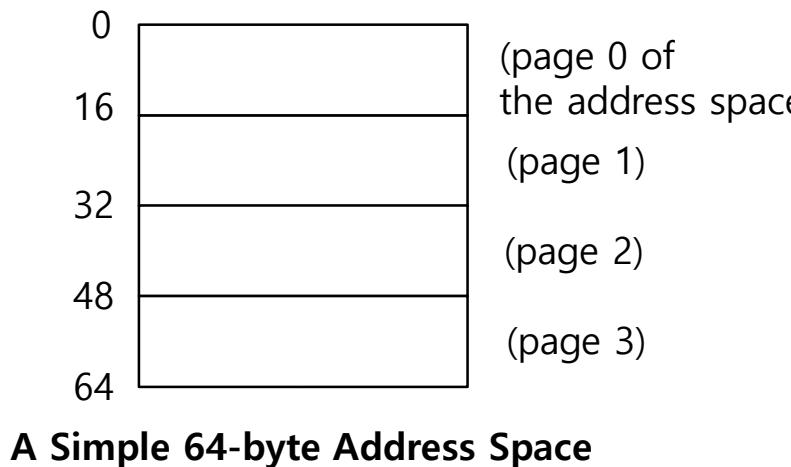
- ▣ Paging **splits up** address space into **fixed-sized** unit called a **page**.
 - ◆ Segmentation: variable size of logical segments(code, stack, heap, etc.)
- ▣ With paging, **physical memory** is also **split** into some number of pages called a **page frame**.
- ▣ **Page table** per process is needed **to translate** the virtual address to physical address.

Advantages Of Paging

- ▣ **Flexibility:** Supporting the abstraction of address space effectively
 - ◆ Don't need assumption how heap and stack grow and are used.
- ▣ **Simplicity:** ease of free-space management
 - ◆ The page in address space and the page frame are the same size.
 - ◆ Easy to allocate and keep a free list

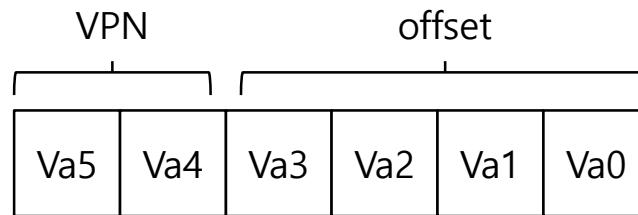
Example: A Simple Paging

- 128-byte physical memory with 16 bytes page frames
- 64-byte address space with 16 bytes pages

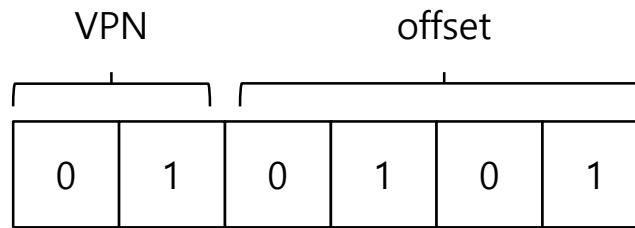


Address Translation

- Two components in the virtual address
 - VPN: virtual page number
 - Offset: offset within the page

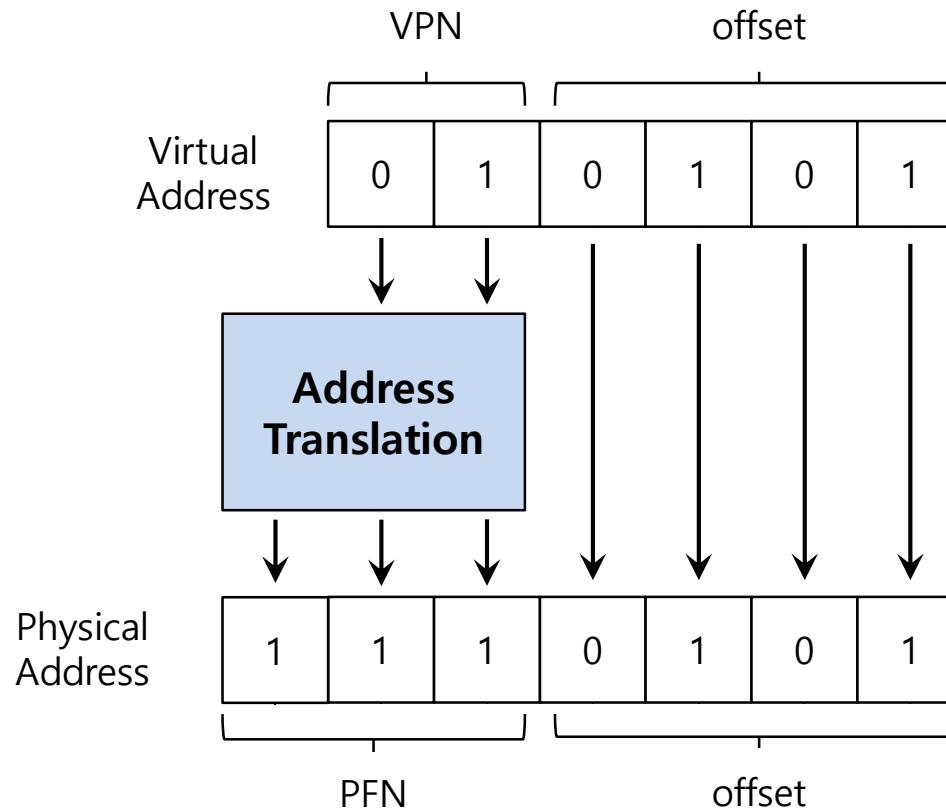


- Example: virtual address 21 in 64-byte address space



Example: Address Translation

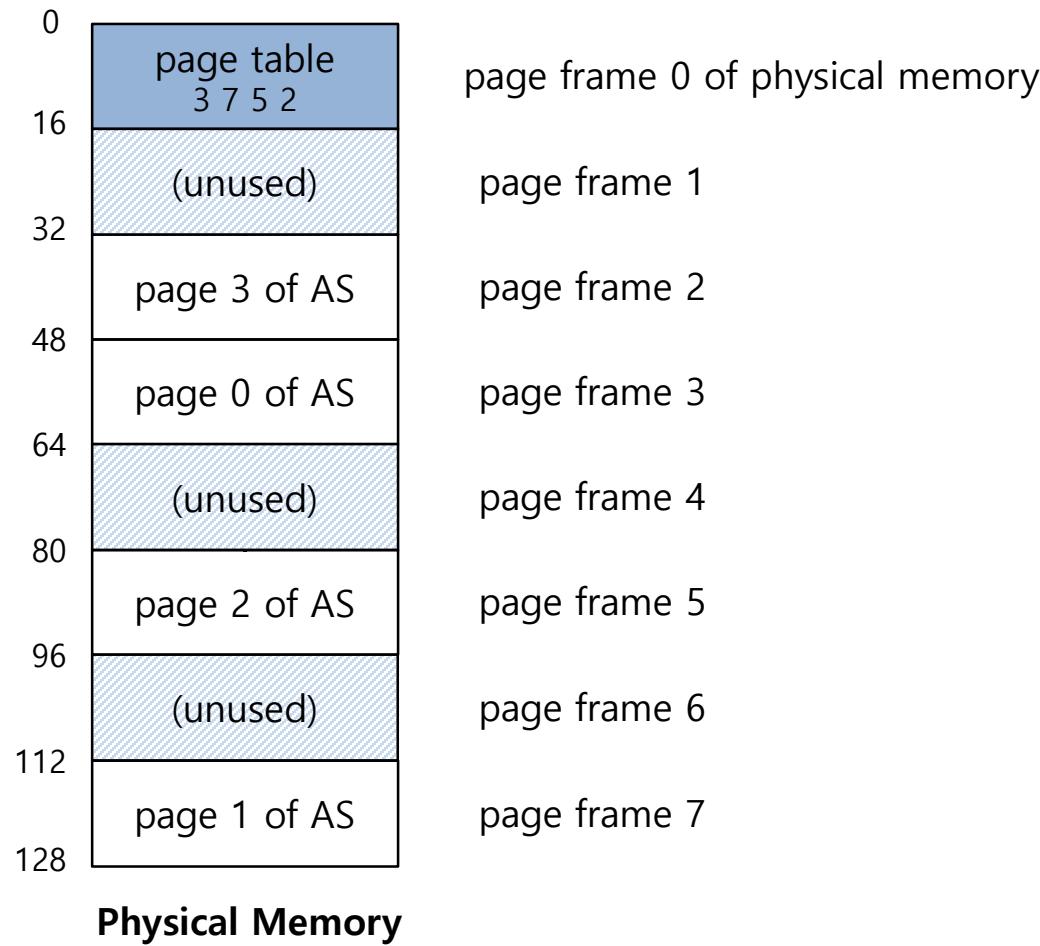
- The virtual address 21 in 64-byte address space



Where Are Page Tables Stored?

- ▣ Page tables can get awfully large.
 - ◆ 32-bit address space with 4-KB pages, 20 bits for VPN
 - Page offset for 4 Kbyte page: 20 bit
 - $4MB = 2^{20} \text{ entries} * 4 \text{ Bytes per page table entry}$
- ▣ Page tables for each process are stored in memory.

Example: Page Table in Kernel Physical Memory



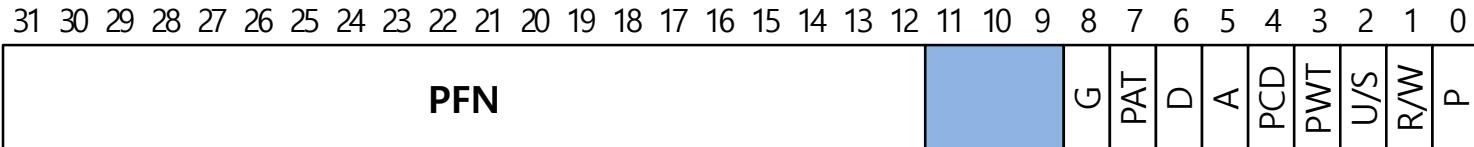
What Is In The Page Table?

- The page table is a **data structure** that is used to map the virtual address to physical address.
 - ◆ Simplest form: a linear page table, an array
- The OS **indexes** the array by VPN, and looks up the page-table entry.

Common Flags Of Page Table Entry

- **Valid Bit:** Indicating whether the particular translation is valid.
- **Protection Bit:** Indicating whether the page could be read from, written to, or executed from
- **Present Bit:** Indicating whether this page is in physical memory or on disk(swapped out)
- **Dirty Bit:** Indicating whether the page has been modified since it was brought into memory
- **Reference Bit(Accessed Bit):** Indicating that a page has been accessed

Example: x86 Page Table Entry



An x86 Page Table Entry(PTE)

- P: present
- R/W: read/write bit
- U/S: supervisor
- A: accessed bit
- D: dirty bit
- PFN: the page frame number

Paging: Too Slow

- ❑ To find a location of the desired PTE, the **starting location** of the page table is **needed**.
- ❑ For every memory reference, paging requires the OS to perform one **extra memory reference**.

Accessing Memory With Paging

```
1 // Extract the VPN from the virtual address  
2 VPN = (VirtualAddress & VPN_MASK) >> SHIFT  
3  
4 // Form the address of the page-table entry (PTE)  
5 PTEAddr = PTBR + (VPN * sizeof(PTE))  
6  
7 // Fetch the PTE  
8 PTE = AccessMemory(PTEAddr)  
9
```

Accessing Memory With Paging

```
10     // Check if process can access the page
11     if (PTE.Valid == False)
12         RaiseException(SEGMENTATION_FAULT)
13     else if (CanAccess(PTE.ProtectBits) == False)
14         RaiseException(PROTECTION_FAULT)
15     else
16         // Access is OK: form physical address and fetch it
17         offset = VirtualAddress & OFFSET_MASK
18         PhysAddr = (PTE.PFN << PFN_SHIFT) | offset
19         Register = AccessMemory(PhysAddr)
```

A Memory Trace

```
int array[1000];  
...  
for (i = 0; i < 1000; i++)  
    array[i] = 0;
```

```
prompt> gcc -o array array.c -Wall -o  
prompt>./array
```

Memory access

```
0x1024 movl $0x0,(%edi,%eax,4) // [edi+eax*4] = 0  
0x1028 incl %eax  
0x102c cmpl $0x03e8,%eax // 0000 0011 1110 10002 = 100010  
0x1030 jne 0x1024
```

A Virtual(And Physical) Memory Trace

