

# **Operating Systems**

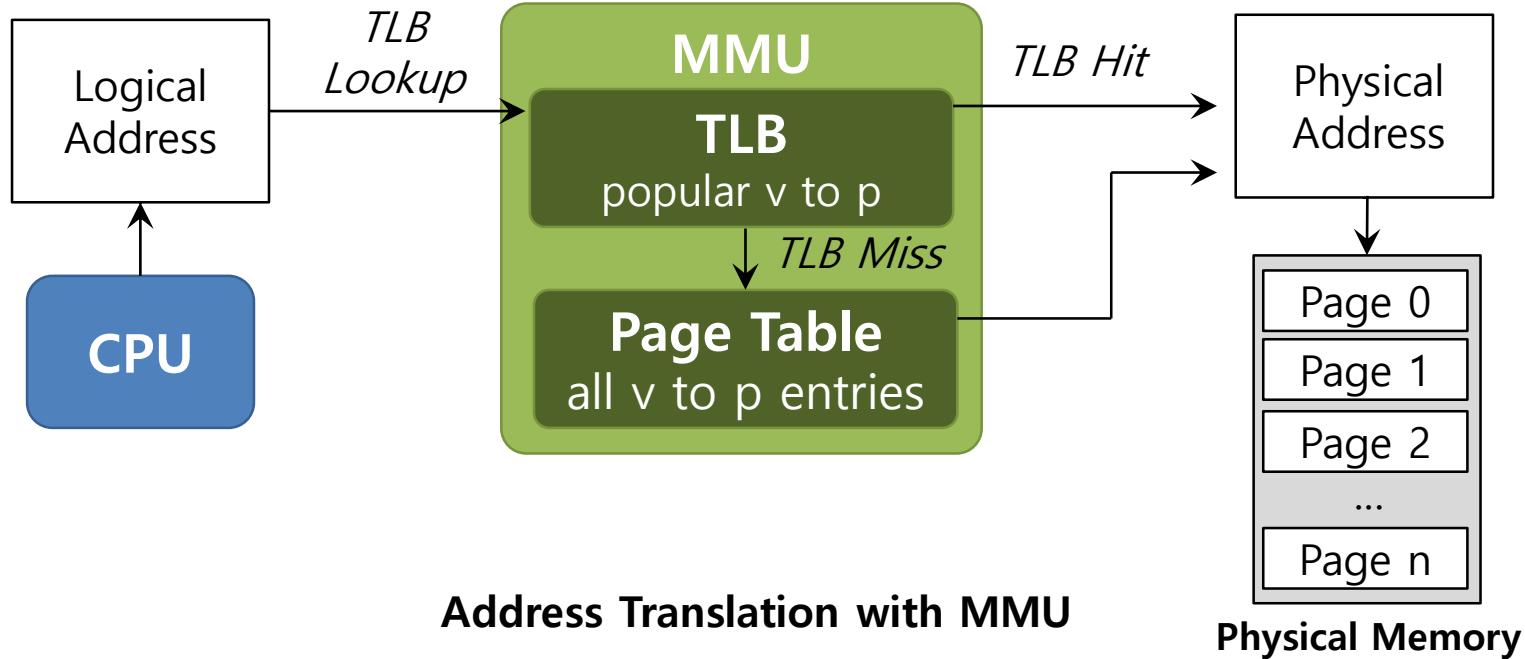
**KAIST**

# 19. Translation Lookaside Buffer

---

# TLB

- Part of the chip's memory-management unit(MMU).
- A hardware cache of **popular** virtual-to-physical address translation.



# TLB Basic Algorithms

```
1: VPN = (VirtualAddress & VPN_MASK ) >> SHIFT  
2: (Success , TlbEntry) = TLB_Lookup(VPN)  
3: if(Success == TRUE) { // TLB Hit  
4:     if(CanAccess(TlbEntry.ProtectBit) == True ) {  
5:         offset = VirtualAddress & OFFSET_MASK  
6:         PhysAddr = (TlbEntry.PFN << SHIFT) | Offset  
7:         AccessMemory( PhysAddr )  
8:     }else RaiseException(PROTECTION_ERROR)
```

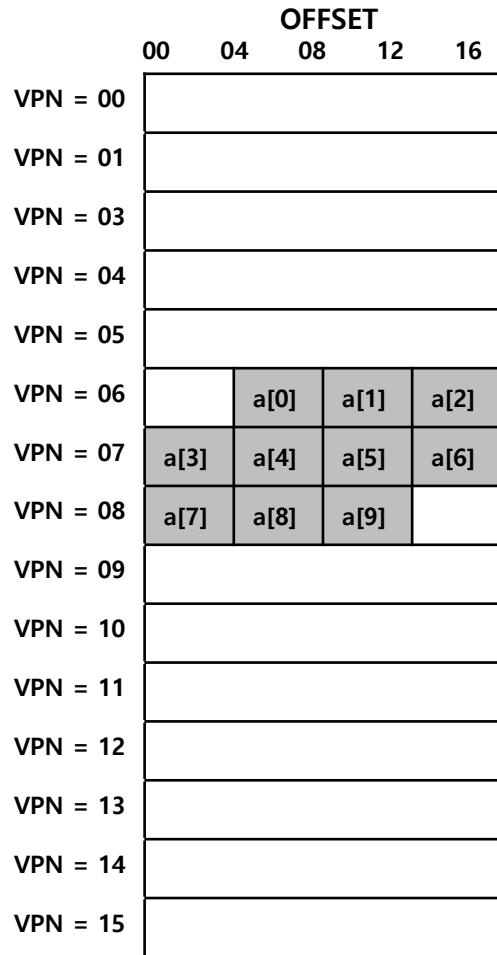
# TLB Basic Algorithms (Cont.)

```
11:      }else{ //TLB Miss  
12:          PTEAddr = PTBR + (VPN * sizeof(PTE))  
13:          PTE = AccessMemory(PTEAddr)  
14:          if(PTE.Valid == False)  
15:              RaiseException SEGFAULT) ;  
16:          else{  
17:              TLB_Insert( VPN , PTE.PFN , PTE.ProtectBits)  
18:              RetryInstruction()  
19:      }
```

- ◆ (11-12 lines) The hardware accesses the page table to find the translation.
- ◆ (16 lines) updates the TLB with the translation.

# Example: Accessing An Array

- How a TLB can improve its performance.



```
0:     int sum = 0 ;  
1:     for( i=0; i<10; i++) {  
2:         sum+=a[i];  
3:     }
```

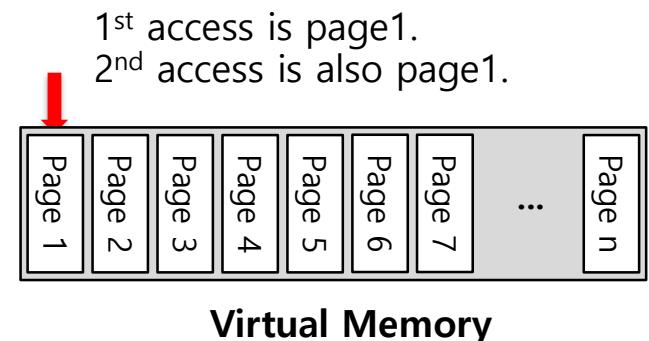
The TLB improves performance  
due to spatial locality

3 misses and 7 hits.  
Thus TLB hit rate is 70%.

# Locality

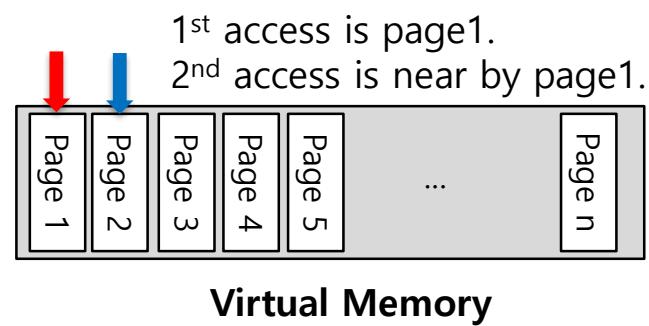
## Temporal Locality

- An instruction or data item that has been recently accessed will likely be re-accessed soon in the future.



## Spatial Locality

- If a program accesses memory at address  $x$ , it will likely soon access memory near  $x$ .



# Who Handles The TLB Miss?

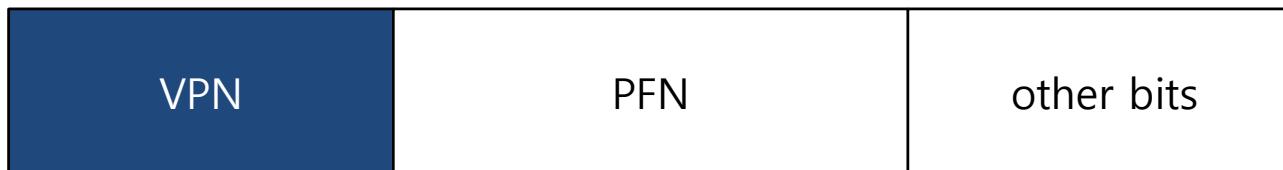
- ▣ Hardware handles the TLB miss entirely on **CISC**.
  - ◆ The hardware has to know exactly where the page tables are located in memory.
  - ◆ The hardware would “walk” the page table, find the correct page-table entry and **extract** the desired translation, **update** and **retry** instruction.
  - ◆ **hardware-managed TLB**.
  - ◆ **Intel x86**
- ▣ **RISC** has what is known as a **software-managed TLB**.
  - ◆ On a TLB miss, the hardware raises exception (trap handler).
    - **Trap handler is code** within the OS that is written with the express purpose of **handling TLB miss**.

# TLB Control Flow algorithm(OS Handled)

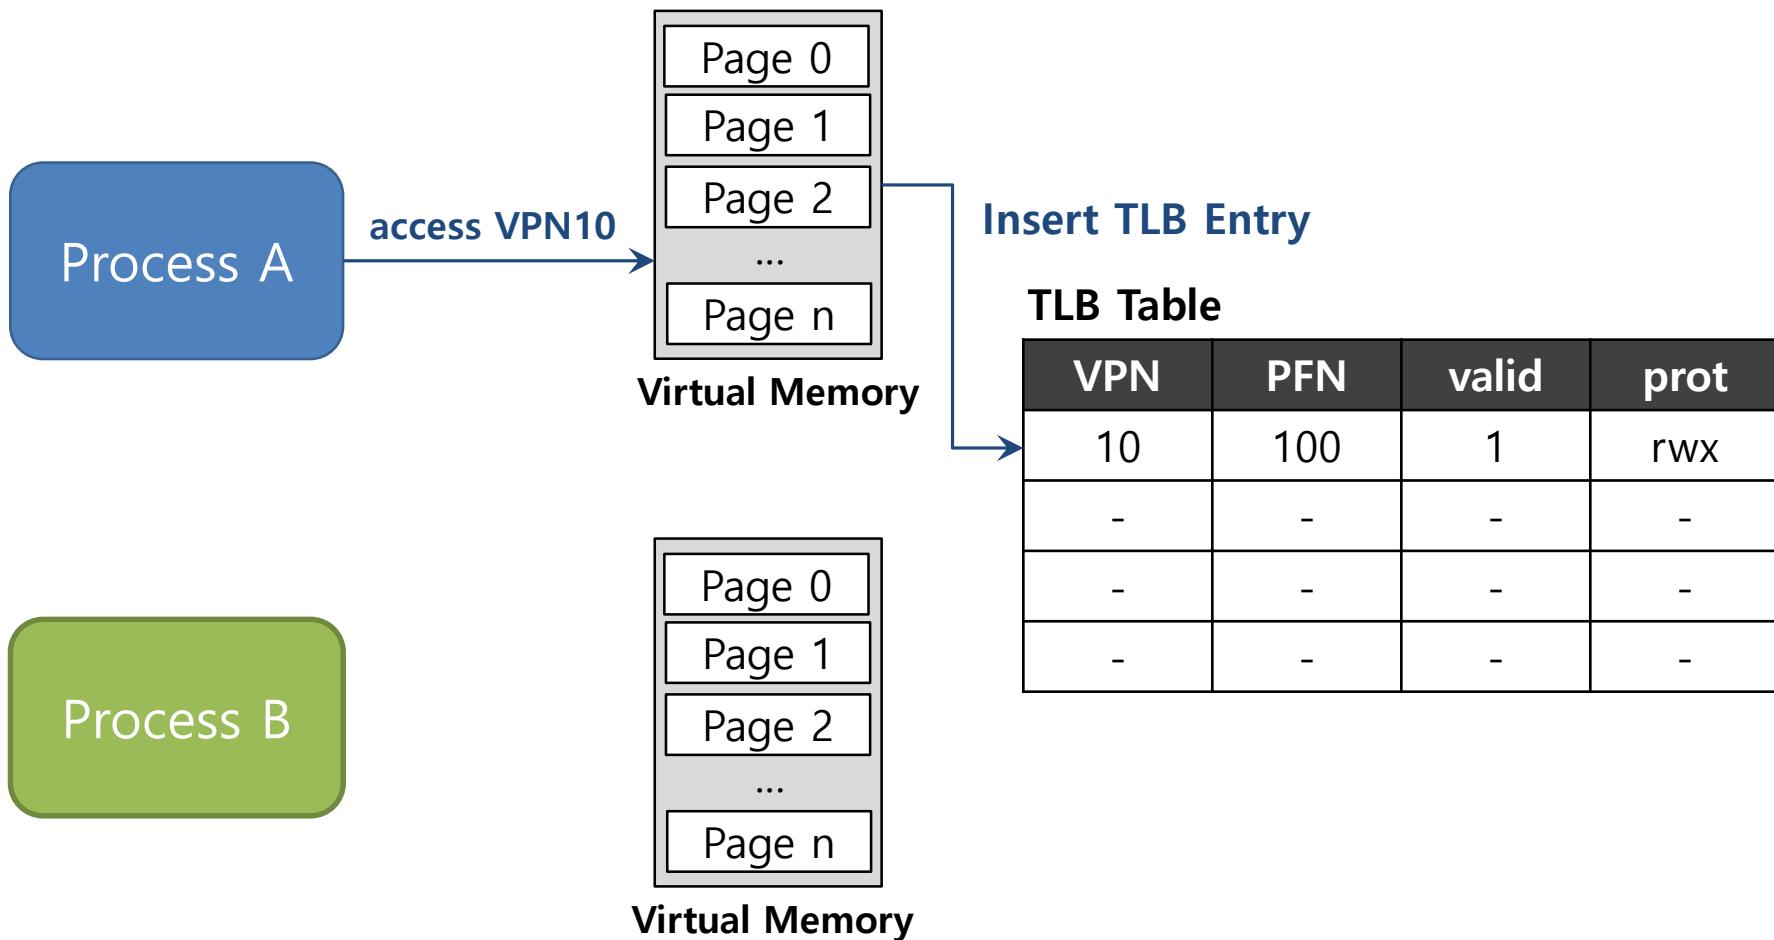
```
1:     VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2:     (Success, TlbEntry) = TLB_Lookup(VPN)
3:     if (Success == True) // TLB Hit
4:         if (CanAccess(TlbEntry.ProtectBits) == True)
5:             Offset = VirtualAddress & OFFSET_MASK
6:             PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7:             Register = AccessMemory(PhysAddr)
8:     else
9:         RaiseException(PROTECTION_FAULT)
10:    else // TLB Miss
11:        RaiseException(TLB_MISS)
```

# TLB entry

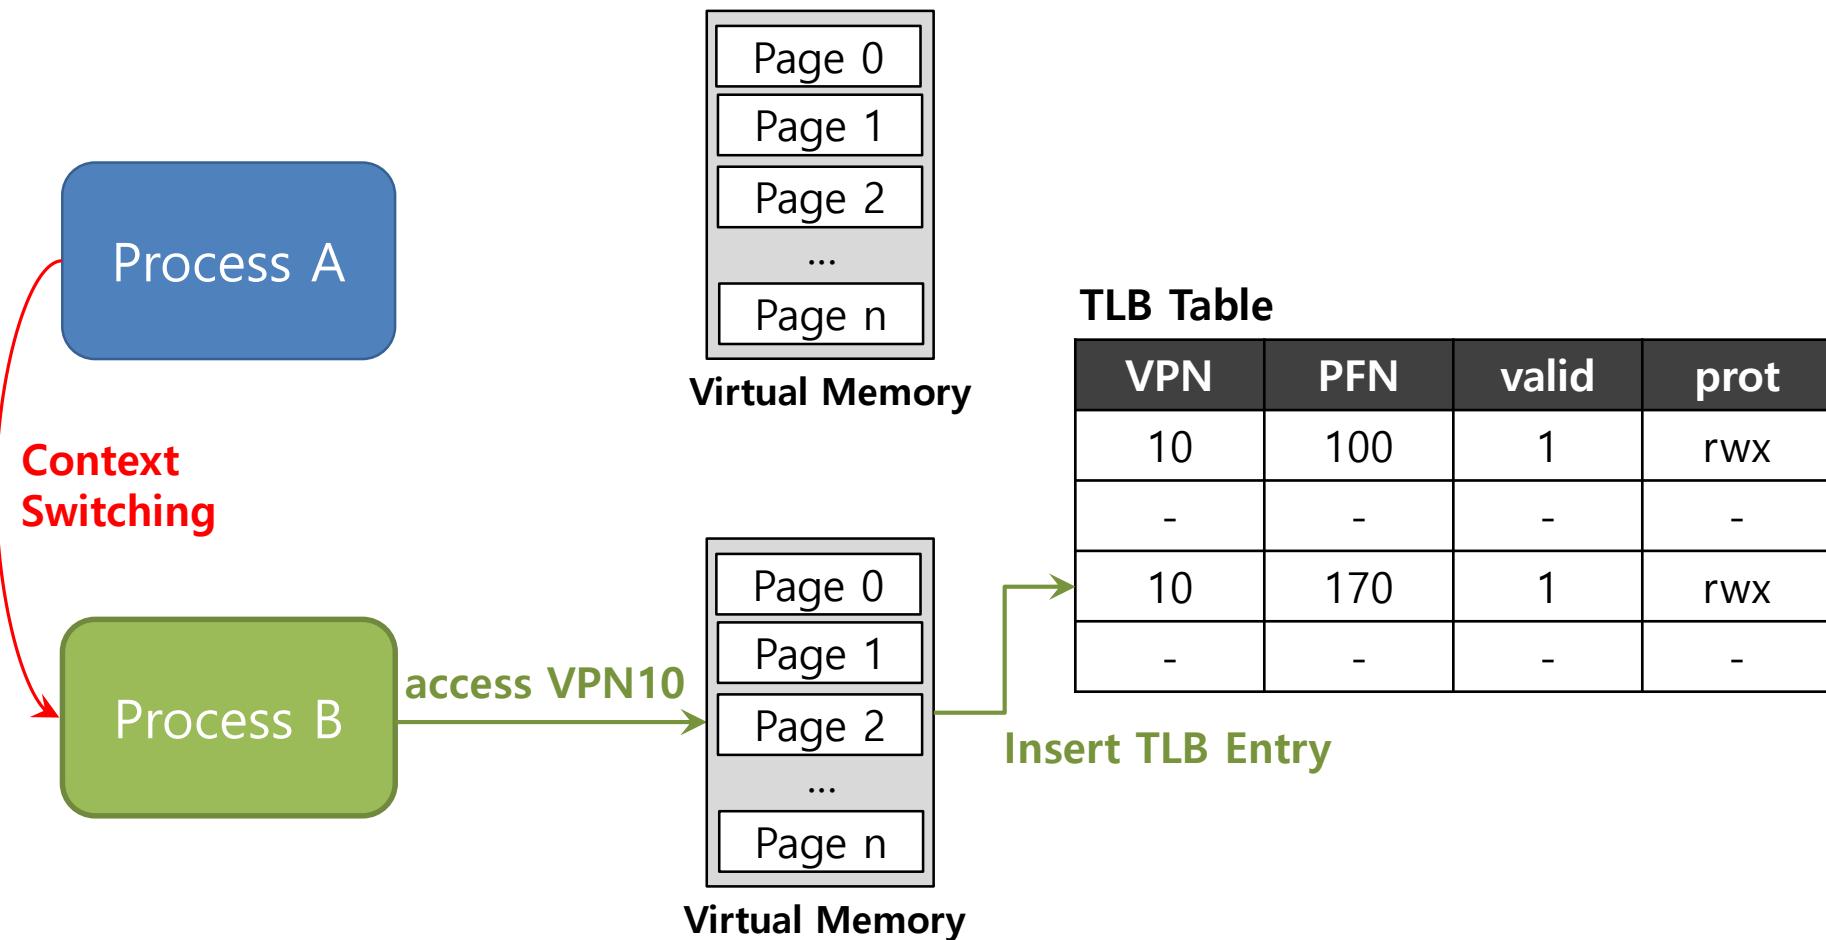
- ❑ TLB is managed by **Full Associative** method.
  - ◆ A typical TLB has 32,64, or 128 entries.
  - ◆ Hardware searches the entire TLB in parallel to find the desired translation.
  - ◆ other bits: valid bits , protection bits, address-space identifier, dirty bit



# TLB Issue: Context Switching

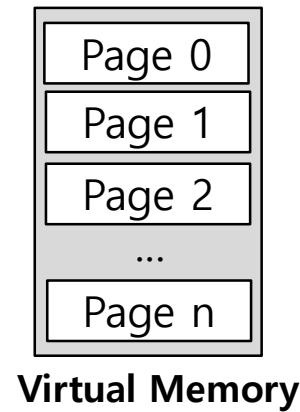


# TLB Issue: Context Switching



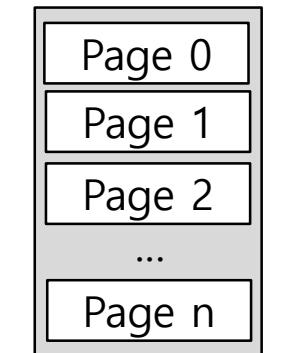
# TLB Issue: Context Switching

Process A



Virtual Memory

Process B



Virtual Memory

TLB Table

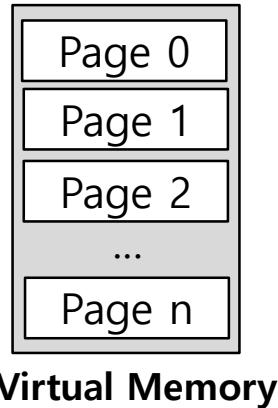
VPN	PFN	valid	prot
10	100	1	rwx
-	-	-	-
10	170	1	rwx
-	-	-	-

Can't Distinguish which entry is meant for which process

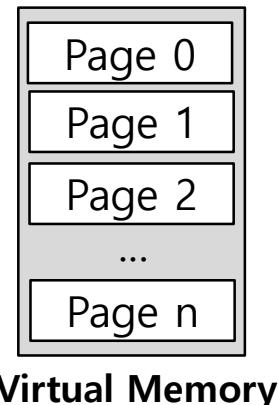
# To Solve Problem

- Provide an address space identifier(ASID) field in the TLB.

Process A



Process B



TLB Table

VPN	PFN	valid	prot	ASID
10	100	1	rwx	1
-	-	-	-	-
10	170	1	rwx	2
-	-	-	-	-

# Another Case

- Two processes share a page.

- Process 1 is sharing physical page 101 with Process2.
- P1 maps this page into the 10<sup>th</sup> page of its address space.
- P2 maps this page to the 50<sup>th</sup> page of its address space.

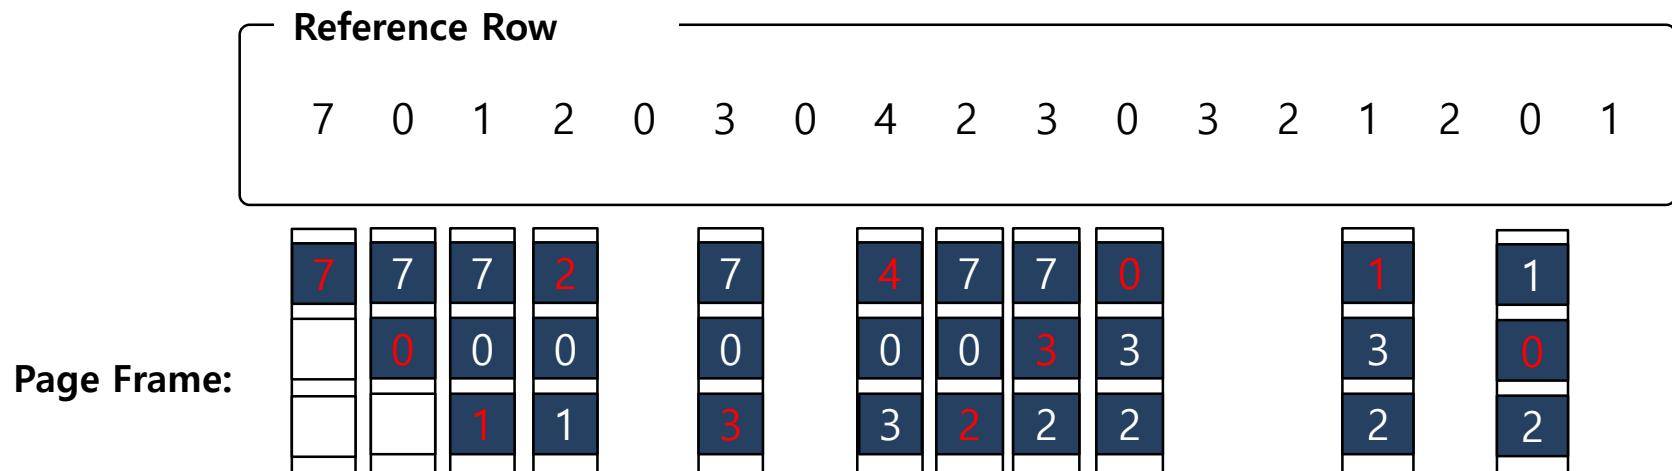
VPN	PFN	valid	prot	ASID
10	101	1	rwx	1
-	-	-	-	-
50	101	1	rwx	2
-	-	-	-	-

Sharing of pages is useful as it reduces the number of physical pages in use.

# TLB Replacement Policy

## □ LRU(Least Recently Used)

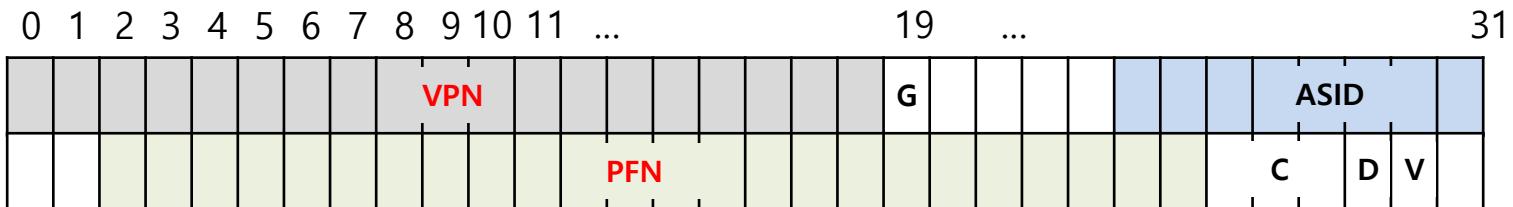
- ◆ Evict an entry that has not recently been used.
- ◆ Take advantage of *locality* in the memory-reference stream.



Total 11 TLB miss

# A Real TLB Entry

All 64 bits of this TLB entry(example of MIPS R4000)



Flag	Content
19-bit VPN	The rest reserved for the kernel.
24-bit PFN	Systems can support up to 64GB of main memory( pages ).
Global bit(G)	Used for pages that are globally-shared among processes.
ASID	OS can use to distinguish between address spaces.
Coherence bit(C)	determine how a page is cached by the hardware.
Dirty bit(D)	marking when the page has been written.
Valid bit(V)	tells the hardware if there is a valid translation present in the entry.