

# Operating Systems



## **22. Swapping: Policies**

---

# Goal of Cache Management

- to minimize the number of cache misses.
- the *average memory access time(AMAT)*.

$$AMAT = (P_{Hit} * T_M) + (P_{Miss} * T_D)$$

Argument	Meaning
$T_M$	The cost of accessing memory
$T_D$	The cost of accessing disk
$P_{Hit}$	The probability of finding the data item in the cache(a hit)
$P_{Miss}$	The probability of not finding the data in the cache(a miss)

# The Optimal Replacement Policy

- ▣ Lead to the fewest number of misses overall.
  - ◆ Replace the page that will be accessed furthest in the future.
  - ◆ Result in the **fewest-possible** cache misses.
- ▣ Serve only as a comparison point, to know how close we are to **perfect**.

# Tracing the Optimal Policy

Reference Row

0 1 2 0 1 3 0 3 1 2 1

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0,1
2	Miss		0,1,2
0	Hit		0,1,2
1	Hit		0,1,2
3	Miss	2	0,1,3
0	Hit		0,1,3
3	Hit		0,1,3
1	Hit		0,1,3
2	Miss	3	0,1,2
1	Hit		0,1,2

Hit rate is  $\frac{Hits}{Hits+Misses} = 54.6\%$

Future is not known.

# A Simple Policy: FIFO

- Pages were placed in a queue when they enter the system.
- When a replacement occurs, the page on the tail of the queue(the **“First-in”** pages) is evicted.
  - ◆ It is simple to implement, but can't determine the importance of blocks.

# Tracing the FIFO Policy

Reference Row

0 1 2 0 1 3 0 3 1 2 1

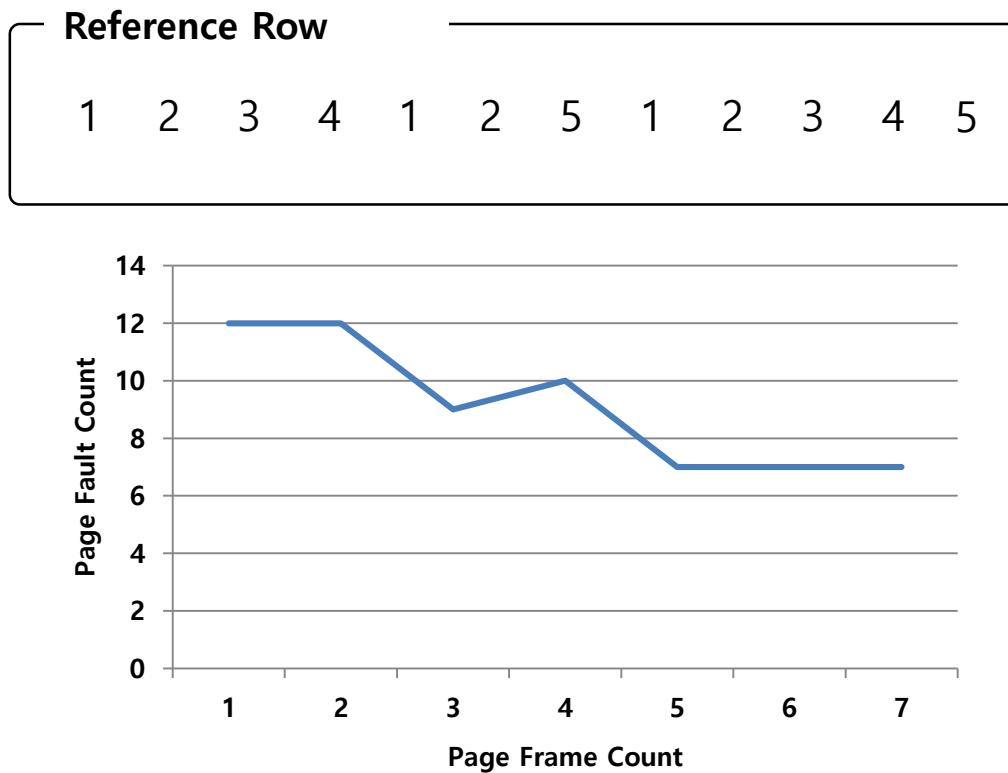
Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0,1
2	Miss		0,1,2
0	Hit		0,1,2
1	Hit		0,1,2
3	Miss	0	1,2,3
0	Miss	1	2,3,0
3	Hit		2,3,0
1	Miss		3,0,1
2	Miss	3	0,1,2
1	Hit		0,1,2

Hit rate is  $\frac{Hits}{Hits+Misses} = 36.4\%$

Even though page 0 had been accessed a number of times, FIFO still kicks it out.

# BELADY' S ANOMALY

- We would expect the cache hit rate to **increase** when the cache gets larger. But in this case, with FIFO, it gets worse.



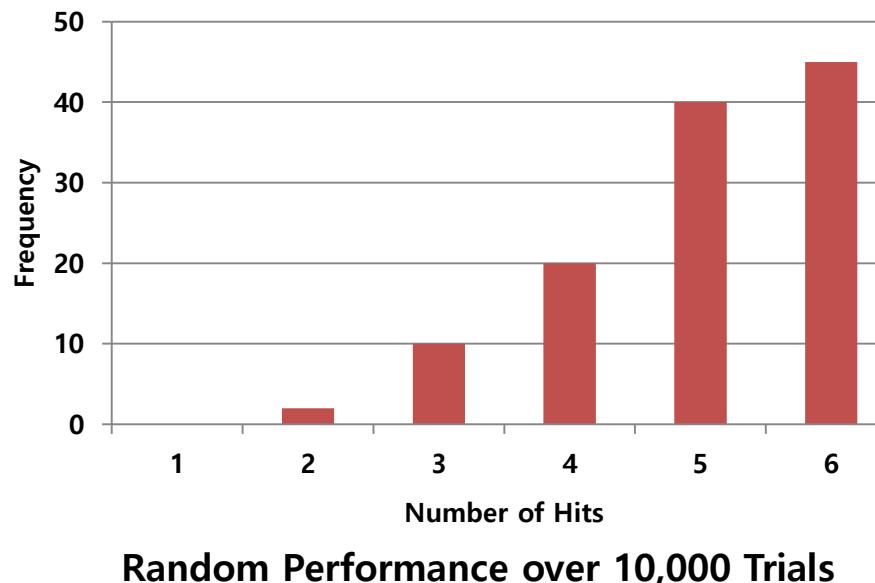
# Another Simple Policy: Random

- ❑ Picks a random page to replace under memory pressure.
  - ◆ It doesn't really try to be too intelligent in picking which blocks to evict.
  - ◆ Random does depends entirely upon how lucky Random gets in its choice.

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0,1
2	Miss		0,1,2
0	Hit		0,1,2
1	Hit		0,1,2
3	Miss	0	1,2,3
0	Miss	1	2,3,0
3	Hit		2,3,0
1	Miss	3	2,0,1
2	Hit		2,0,1
1	Hit		2,0,1

# Random Performance

- ❑ Sometimes, Random is as good as optimal, achieving 6 hits on the example trace.



# Using History

- Learn on the past and use **history**.
  - Two type of historical information.

Historical Information	Meaning	Algorithms
<b>recency</b>	The more recently a page has been accessed, the more likely it will be accessed again	LRU
<b>frequency</b>	If a page has been accessed many times, It should not be replaced as it clearly has some value	LFU

# Using History : LRU

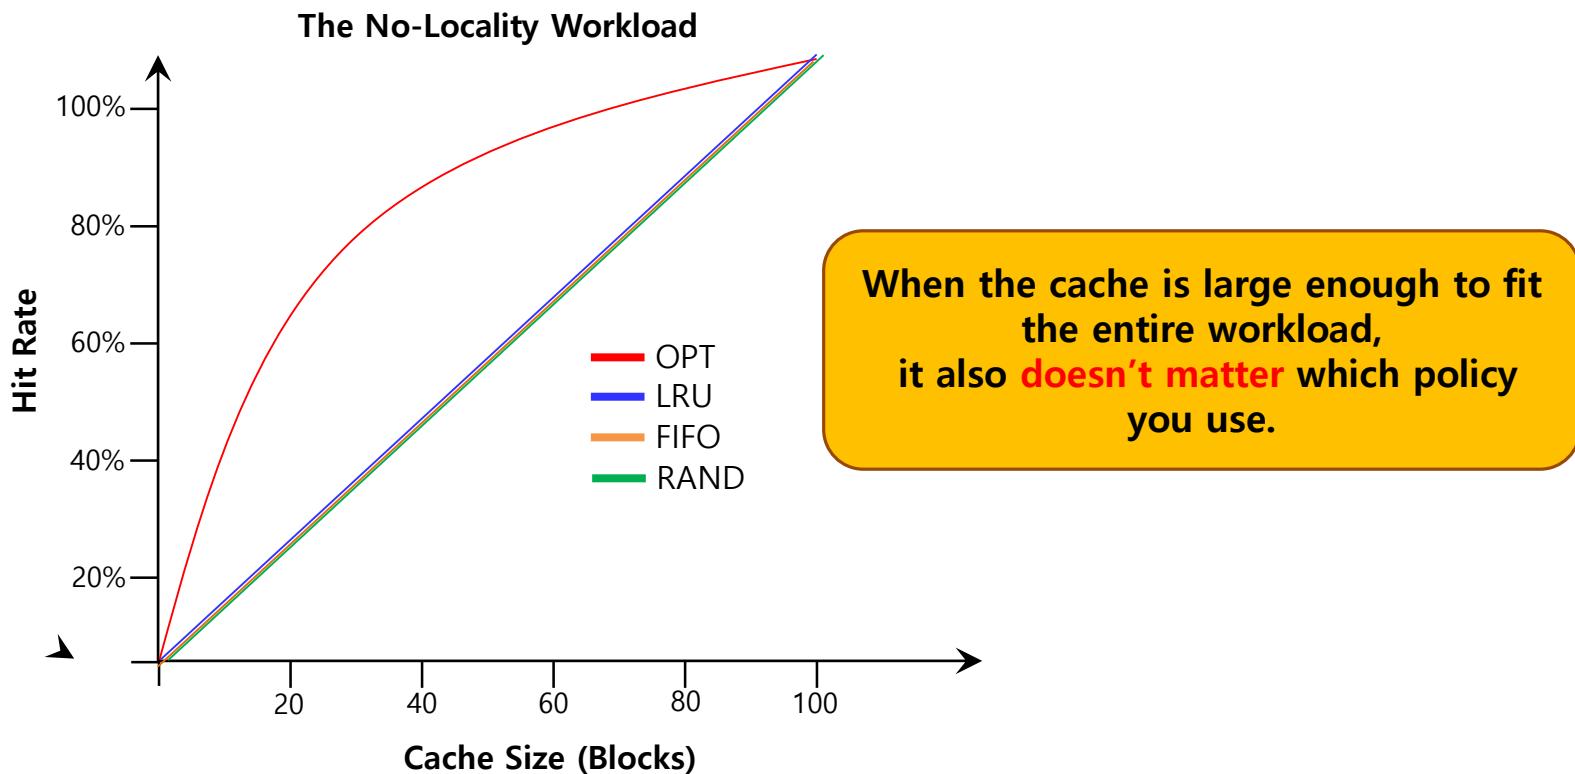
- Replace the least-recently-used page.

Reference Row										
0	1	2	0	1	3	0	3	1	2	1

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0,1
2	Miss		0,1,2
0	Hit		1,2,0
1	Hit		2,0,1
3	Miss	2	0,1,3
0	Hit		1,3,0
3	Hit		1,0,3
1	Hit		0,3,1
2	Miss	0	3,1,2
1	Hit		3,2,1

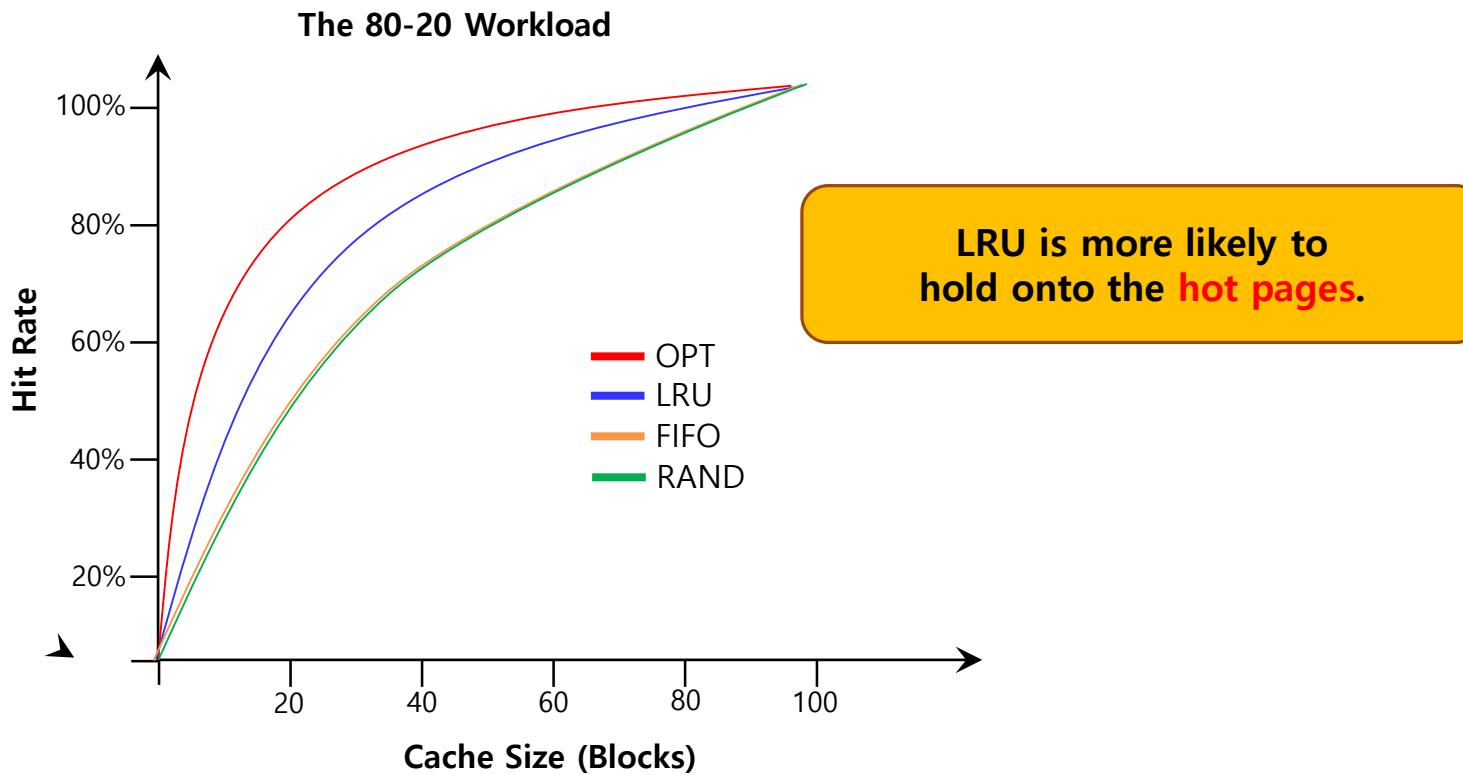
# Workload Example : The No-Locality Workload

- Each reference is to a random page within the set of accessed pages.
  - Workload accesses 100 unique pages over time.
  - Choosing the next page to refer to at random



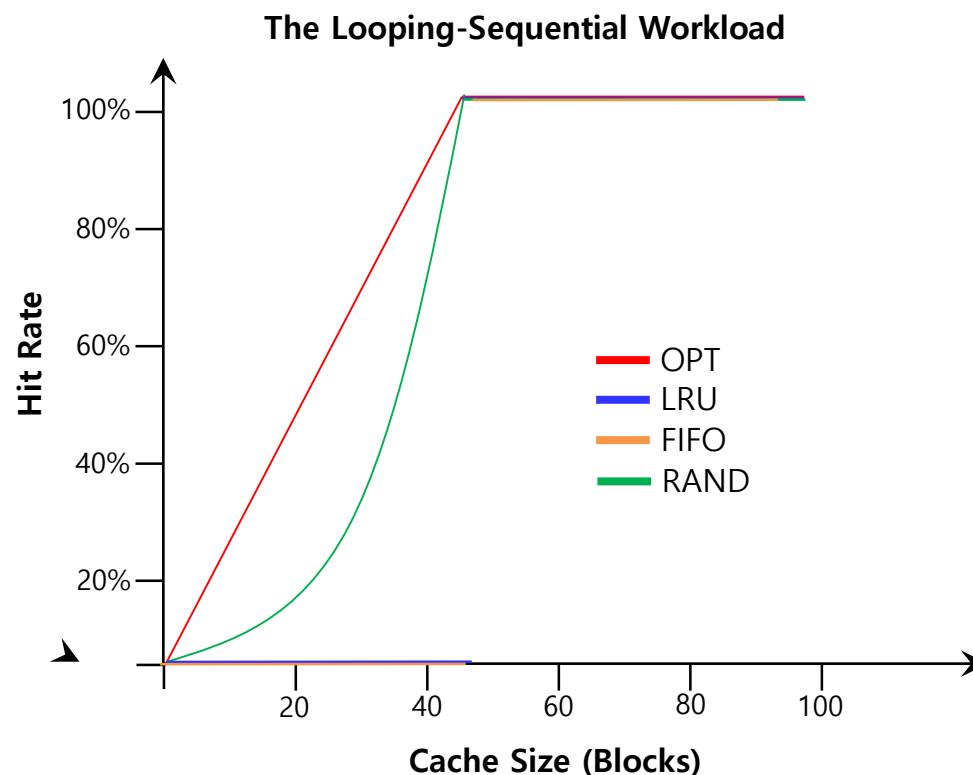
# Workload Example : The 80-20 Workload

- Exhibits locality: 80% of the **reference** are made to 20% of the page
- The remaining 20% of the **reference** are made to the remaining 80% of the pages.



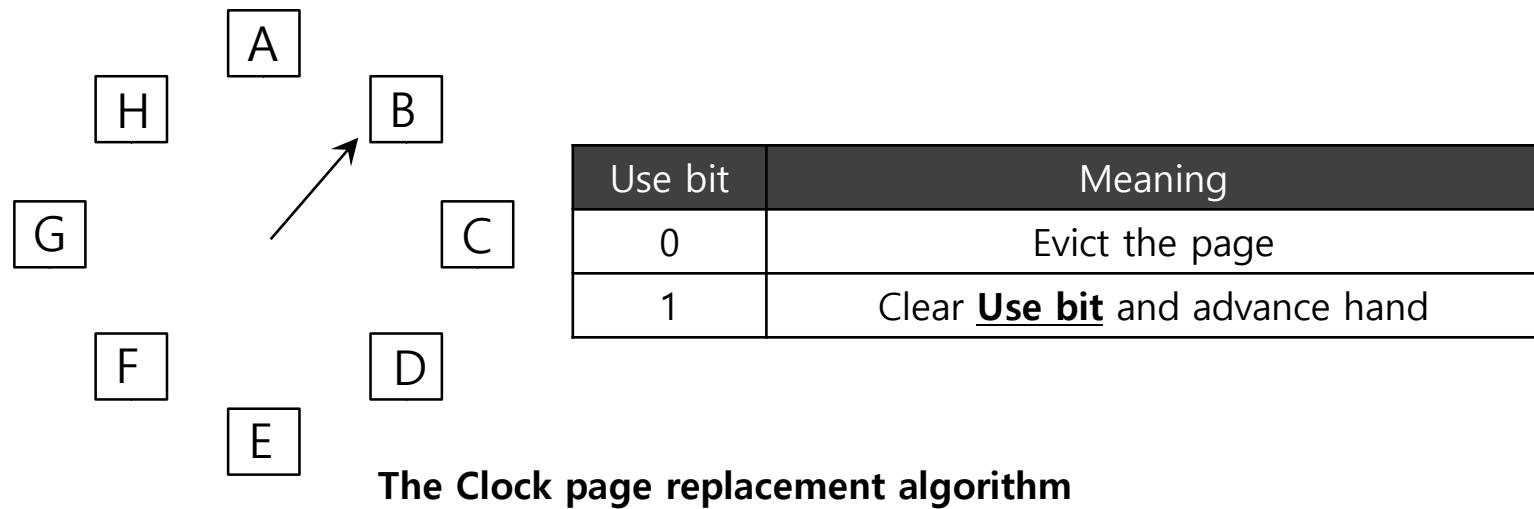
# Workload Example : The Looping Sequential

- Refer to 50 pages in sequence.
  - Starting at 0, then 1, ... up to page 49, and then we Loop, repeating those accesses, for total of 10,000 accesses to 50 unique pages.



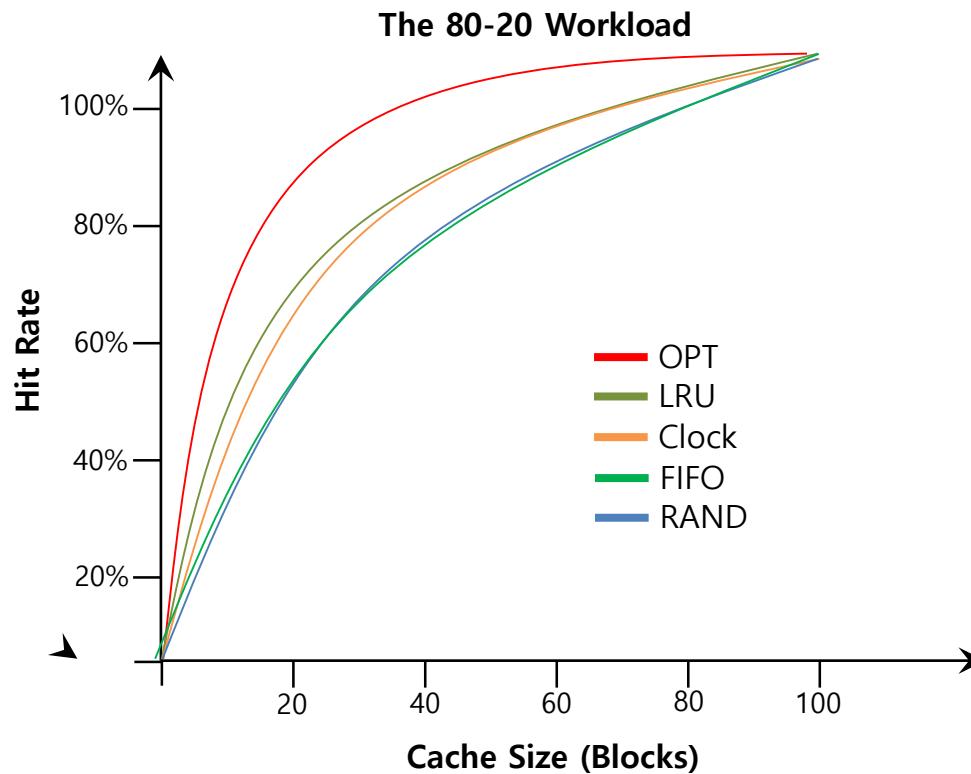
# Approximating LRU: Clock Algorithm

- Require hardware support: a use bit
  - ◆ Whenever a **page is referenced**, the use bit is set by hardware to 1.
  - ◆ Hardware **never** clears the bit, though; that is the responsibility of the OS
- Clock Algorithm
  - ◆ All pages of the system arranges in a circular list.
  - ◆ A clock hand points to some particular page to begin with.
  - ◆ The algorithm continues until it finds a use bit that is set to 0.



# Workload with Clock Algorithm

- Clock algorithm doesn't do as well as perfect LRU, it does better then approach that don't consider history at all.

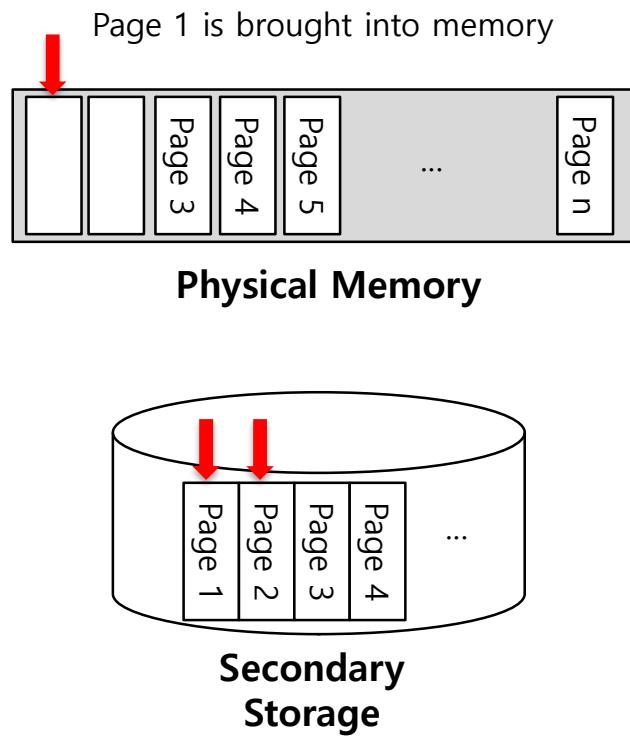


# Considering Dirty Pages

- ▣ The hardware includes a modified bit (a.k.a dirty bit)
  - ◆ Page has been modified and is thus dirty, it must be written back to disk to evict it.
  - ◆ Page has not been modified, the eviction is free.

# Prefetching

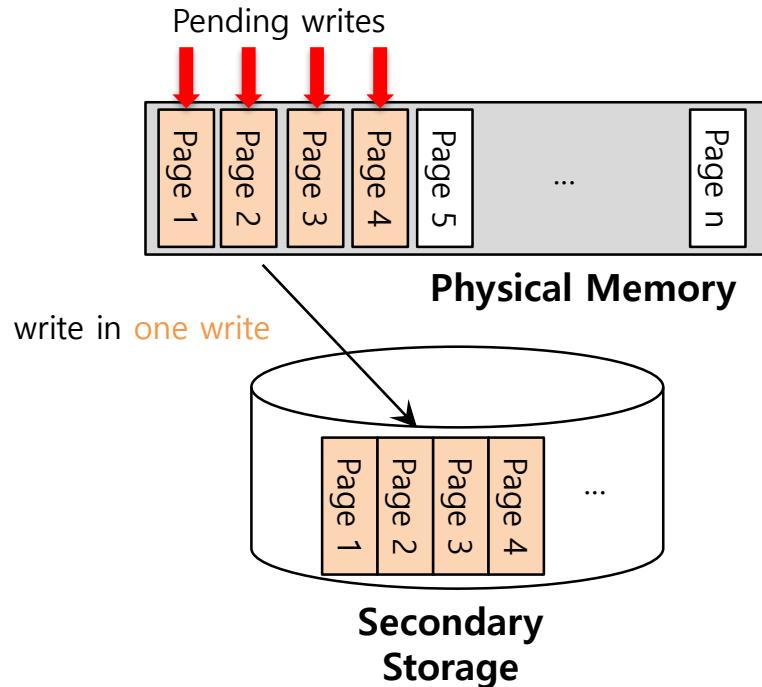
- The OS guesses that a page is about to be used, and thus bring it in ahead of time.



Page 2 likely **soon be accessed** and  
thus should be brought into memory too

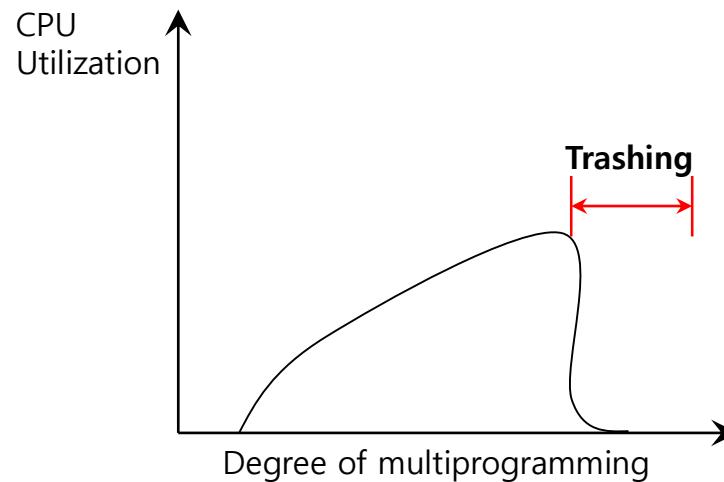
# Clustering, Grouping

- Collect a number of **pending writes** together in memory and write them to disk in **one write**.
  - Perform a **single large write** more efficiently than **many small ones**.



# Thrashing

- Memory is **oversubscribed** and the memory demands of the set of running processes **exceeds** the available physical memory.
  - Decide not to run a subset of processes.
  - Reduced set of processes working sets fit in memory.



# Summary

- ▣ Swapping: use part of disk as memory
- ▣ LRU, LFU, RANDOM, FIFO
- ▣ Approximation to LRU: Clock
- ▣ Making the disk IO in larger unit
  - ◆ Clustering
  - ◆ Grouping
  - ◆ prefetching