

Operating Systems



26. Concurrency: An Introduction

Thread


- ▣ A new abstraction for a single running process
- ▣ Multi-threaded program
 - ◆ A multi-threaded program has more than one point of execution.
 - ◆ Multiple PCs (Program Counter)
 - ◆ They **share** the share the same **address space**.

Context switch between threads

- ▣ Each thread has its own program counter and set of registers.
 - ◆ One or more **thread control blocks(TCBs)** are needed to store the state of each thread.
- ▣ When switching from running one (T1) to running the other (T2),
 - ◆ The register state of T1 be saved.
 - ◆ The register state of T2 restored.
 - ◆ The **address space remains** the same.

The stack of the relevant thread

- There will be one stack per thread.




A Single-Threaded Address Space

The code segment:
where instructions live

The heap segment:
contains malloc'd data
dynamic data structures
(it grows downward)

(it grows upward)
The stack segment:
contains local variables
arguments to routines,
return values, etc.



Two threaded Address Space

Why Use Threads?

- ▣ Parallelism
 - ◆ Single-threaded program: the task is straightforward, but slow.
 - ◆ Multi-threaded program: natural and typical way to make programs run faster on modern hardware.
 - ◆ **Parallelization:** The task of transforming standard **single-threaded** program into a program that does this sort of work on multiple CPUs.
- ▣ Avoid blocking program progress due to slow I/O.
 - ◆ Threading enables **overlap** of I/O with other activities within a single program.
 - ◆ It is much like **multiprogramming** did for processes across programs.

An Example: Thread Creation

Simple Thread Creation Code (t0.c)

```
1 #include <stdio.h>
2 #include <assert.h>
3 #include <pthread.h>
4 #include "common.h"
5 #include "common_threads.h"
6
7 void *mythread (void *arg) {
8     printf ("%s\n", (char *) arg);
9     return NULL;
10 }
11
12 int main (int argc, char *argv[]) {
13     pthread_t p1, p2;
14     int rc;
15     printf("main: begin\n");
16     Pthread_create(&p1, NULL, mythread, "A");
17     Pthread_create(&p2, NULL, mythread, "B");
18     // join waits for three threads to finish
19     Pthread_join(p1, NULL);
20     Pthread_join(p2, NULL);
21     printf("main: end\n");
22     return 0;
23 }
```

Thread Trace (1)

main	Thread 1	Thread 2
starts running		
prints "main: begin"		
creates Thread 1		
creates Thread 2		
waits for T1		
	runs	
	prints "A"	
	returns	
waits for T2		
		runs
		prints "B"
		returns
prints "main: end"		

Thread Trace (2)

main	Thread 1	Thread 2
starts running		
prints "main: begin"		
creates Thread 1		
	runs	
	prints "A"	
	returns	
creates Thread 2		
	runs	
	prints "B"	
	returns	
waits for T1		
returns immediately; T1 is done		
waits for T2		
returns immediately; T2 is done		
prints "main: end"		

Thread Trace (3)

main	Thread 1	Thread 2
starts running		
prints "main: begin"		
creates Thread 1		
creates Thread 2		
		runs
		prints "B"
		returns
waits for T1		
	runs	
	prints "A"	
	returns	
waits for T2		
returns immediately; T2 is done		
prints "main: end"		

Race condition

- Increasing a value of a variable

- counter = counter + 1

```
105      mov 0x8049a1c, %eax
108      add $0x1, %eax
113      mov %eax, 0x8049a1c
```



Race condition

Example with two threads

- counter = counter + 1 (default is 50)
- We expect the result is 52. However,

OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
		<i>before critical section</i>			
		mov 0x8049a1c, %eax	100	0	50
		add \$0x1, %eax	105	50	50
			108	51	50
interrupt					
	save T1's state				
	restore T2's state		100	0	50
		mov 0x8049a1c, %eax	105	50	50
		add \$0x1, %eax	108	51	50
		mov %eax, 0x8049a1c	113	51	51
interrupt					
	save T2's state				
	restore T1's state		108	51	50
		mov %eax, 0x8049a1c	113	51	51

A few terminologies

- ▣ Race condition:
 - ◆ the results depend on the timing execution of the code.
 - ◆ Result is indeterminate.
- ▣ Critical section
 - ◆ A piece of code that **accesses a shared variable** and must not be concurrently executed by more than one thread.
 - ◆ Multiple threads executing critical section can result in a race condition.
 - ◆ Need to support **atomicity** for critical sections (**mutual exclusion**)



The wish for atomicity

- ☐ Ideal approach; make the increment as a single assembly instruction

```
memory-add 0x8049alc, $0x1
```

- ☐ In general, we do not have such instruction. Instead, we use lock.
- ☐ Ensure that any such critical section executes as if it were a single atomic instruction (**execute a series of instructions atomically**).

```
1  lock_t mutex;
2  . . .
3  lock(&mutex);
4  balance = balance + 1; → Critical section
5  unlock(&mutex);
```